

Response to Reviewer Comments

Revised Submission

We sincerely thank all three reviewers and the meta-reviewer for their careful, constructive, and detailed feedback. We have substantially revised the manuscript to address each revision request. In this response document, we summarize **what changes were made, the associated revision actions taken to address each concern, and our response and discussion**. All changes in the revised manuscript are highlighted in blue for ease of inspection.

To make the response easy to navigate, each item begins with the corresponding reviewer or meta-reviewer comment, followed by a concise summary of the revision actions and an explanation of how these changes address the comment.

RESPONSE TO THE META-REVIEWER

MR.1 — Strengthening the theoretical framework

(addresses Rev1.O2/O3, Rev1.D2, Rev2.D1/D2, Rev3.D2/D3)

What we changed:

- We rewrote the proof of Theorem 1 into three explicit cases: (1) the uniform case with $\alpha = 1$, (2) the uniform case with general $\alpha > 0$, and (3) the bounded non-uniform density case. [Sec. 4.5, proof of Theorem 1]
- We added a detailed verification of the Differential Equation Method (DEM) conditions, including Boundedness, Trend, and Lipschitz continuity. [Sec. 4.5, proof of Theorem 1]
- We revised the proof of Lemma 3 to clarify that the Binomial model is used through a conditional independence argument, conditional on the currently selected nearest neighbor p_t^* . [Sec. 4.3, proof of Lemma 3, first paragraph]
- We added a short “Preview of the main result” paragraph at the beginning of Section 4, so that Theorem 1 is stated before the supporting lemmas. [Sec. 4, second paragraph]
- We split the former Section 4.3 into two subsections: Section 4.3, “Two-Dimensional Case” (Lemma 3 only), and Section 4.4, “High-Dimensional Settings” (Lemma 4). [Sec. 4.3; Sec. 4.4]
- We rewrote Section 6.2 to derive $R^* \propto \log n / \alpha^2$ step by step from the construction-cost model. [Sec. 6.2, second paragraph]
- We added an explanation of why the analytically derived R^* can slightly improve Recall by avoiding locally suboptimal truncation choices. [Sec. 7.3, third paragraph, search-performance discussion]
- We corrected the missing-reference marker “??” in the proof of Lemma 4 and added a footnote clarifying the measure-theoretic meaning of “almost surely.” [Sec. 4.4, proof of Lemma 4]
- We revised the statement of Theorem 2 to make the uniform-distribution assumption explicit and added a discussion of its extension to non-uniform distributions as future work. [Sec. 5 Theorem 2 statement; Sec. 9, newly added third paragraph in the Conclusions]

The meta-review highlights an important theoretical concern: the gap between the broadly stated Theorem 1 and the narrower assumptions used in the supporting lemmas. We have revised the manuscript to make this extension argument explicit.

Closing the gap between Theorem 1 and the supporting lemmas. The original proof developed the main argument first under the uniform-density model with $\alpha = 1$, but the extension to more general settings was not stated explicitly enough. In the revised manuscript, the proof of Theorem 1 is organized into three explicit cases. Case 1 treats the uniform $\alpha = 1$ setting as the geometric base case. Case 2 extends the argument to general fixed $\alpha > 0$ using the corresponding Apollonius blocking region, which reduces to the spherical-cap geometry when $\alpha = 1$. Case 3 further extends the result to bounded non-uniform densities satisfying $0 < m \leq f \leq M < \infty$. This condition captures many locally regular non-uniform settings while excluding degenerate cases with vanishing density or singular concentration. Let μ denote the induced non-uniform candidate distribution, and let $\pi_t(\mu \mid p_t^*)$ denote the conditional probability that a remaining candidate is pruned at iteration t , given the currently selected nearest neighbor p_t^* , under distribution μ . Under this assumption, we prove

$$\pi_t(\mu \mid p_t^*) \geq \frac{m}{M} c(d, \alpha) =: c_* > 0.$$

The key observation is that the subsequent two-phase pruning analysis (Chernoff bound, Borel–Cantelli argument, and Wormald’s DEM) requires only a uniform positive lower bound on the per-iteration pruning probability. Once $c_* > 0$ is established, the rest of the proof proceeds with $c(d, 1)$ replaced by c_* . The dimension d enters through the pruning-probability estimate and is absorbed into the constants $c(d, 1)$ and $c(d, \alpha)$, so the argument remains valid in general d -dimensional settings.

MR.2 — Improving the presentation

(addresses Rev1.R1)

What we changed: We split the former Section 4.3, which was titled “Two-Dimensional Case” but contained both Lemma 3 in two dimensions and Lemma 4 in d dimensions, into two subsections. The new **Section 4.3** contains only Lemma 3, while the new **Section 4.4, “High-Dimensional Settings”** contains Lemma 4. [Sec. 4.3; Sec. 4.4]

The reviewer correctly identified a mismatch between the subsection title and its contents. Our original intention was to use the low-dimensional case as an illustrative starting point: in two dimensions, the pruning acceleration effect can be characterized geometrically, which is formalized in Lemma 3. In the high-dimensional setting, the same phenomenon is captured as *sublinear-time progress*, namely that a linear number of points can be processed within a sublinear number of iterations; this is formalized in Lemma 4. The revised organization separates these two results, so each subsection title now accurately reflects the material it contains.

MR.3 — Expanding the experimental evaluation

(addresses Rev2.D3, Rev3.D1, Rev3.W3)

What we changed:

- We added a new subsection, 7.5 “**Ablation: Parameter Reuse Across Builds**,” together with a new figure reporting the parameter-reuse results. In this ablation, the baseline sweeps R once and then reuses the selected value across subsequent builds, so that no additional tuning cost is charged to the baseline. [Sec. 7.5]
- We added four additional baselines covering major families of recent ANNS methods: **VSAG**, **SymphonyQG**, **RaBitQ**, and **DET-LSH**. We also added Recall-QPS and construction-time comparisons against these baselines, together with the corresponding bibliography entries. [Sec. 7.1.6 “Additional baselines”; Sec. 7.3, fifth new paragraph “Comparison with Additional Baselines”]
- We added a billion-scale evaluation on three standard benchmarks: SIFT1B, DEEP1B, and SpaceV1B, following the protocol of recent large-scale ANNS evaluations such as TriBase. [New Sec. 7.6, billion-scale block, third–fourth paragraphs; Figure 5]

Parameter-reuse ablation (Rev2.D3). In the new ablation, the baseline first performs a parameter sweep over R on one dataset and then reuses the selected value on subsequent builds from the same dataset family. Thus, the reuse baseline pays no additional tuning cost on the new builds. Even under this favorable setting for the baseline, OPT-SNG still achieves average construction-time reductions of $1.09\times$ for DiskANN-Vamana, $1.65\times$ for NSG, and $1.74\times$ for HNSW. Although the speedup becomes smaller after amortizing the tuning cost, OPT-SNG still provides direct construction-time improvements. This confirms that the reported gains are not merely due to parameter-sweep overhead; they also stem from the improved quality of the analytically selected R^* , which remains beneficial under parameter reuse.

New baselines (Rev3.D1). We added four further baselines, covering graph-based methods (VSAG and SymphonyQG), a quantization-based method (RaBitQ), and a hashing-based method (DET-LSH). The revised comparison shows that OPT-SNG matches or surpasses VSAG and SymphonyQG along the Recall-QPS frontier on all datasets, while building the index about $1.2\times$ faster than VSAG and $1.8\times$ faster than SymphonyQG on average. Compared with DET-LSH and RaBitQ, OPT-SNG achieves better construction and search efficiency at comparable recall.

Billion-scale evaluation (Rev3.W3). Adding billion-scale datasets makes the empirical evaluation more complete and provides a more direct test of scalability beyond the million-scale setting. Accordingly, we added a new evaluation block on SIFT1B, DEEP1B, and SpaceV1B. Across all three datasets, OPT-SNG preserves the construction-time advantage observed at the million-point scale. In particular, OPT-Vamana achieves graph-construction speedups of $2.48\times$ on SIFT1B, $1.03\times$ on DEEP1B, and $1.07\times$ on SpaceV1B. Moreover, at fixed latency targets, OPT-SNG matches or improves Recall@10 compared with the original version on all three datasets. Due to hardware constraints, the absolute QPS in our billion-scale experiments is lower than what may be achieved on large-scale production systems. However, all methods are evaluated under the same hardware setting, so the comparison still reflects the relative effectiveness of the proposed analytical parameter selection.

RESPONSE TO REVIEWER #1

We thank the reviewer for the careful reading and helpful comments. We address each point below.

O1 / R1 – Section 4.3 title mismatch

O1. I have some confusions about Section 4.3. The title of Section 4.3 is “Two-Dimensional Case”. It contains two lemmas: Lemma 3 for the two-dimensional space $\mathbb{R}^2$, and Lemma 4 for the d -dimensional space $\mathbb{R}^d$. Should the title of Section 4.3 be revised OR Lemma 4 be moved to another subsection?

R1. Examine whether the title of Section 4.3 matches well with the lemmas in Section 4.3.

What we changed: We split the original Section 4.3 into two subsections so that each subsection title accurately reflects its content. Specifically, Lemma 3 remains in Section 4.3, now titled “Two-Dimensional Case,” while Lemma 4 has been moved to the new Section 4.4, titled “High-Dimensional Settings.” [Sec. 4.3 “Two-Dimensional Case”; Sec. 4.4 “High-Dimensional Settings”]

Response and discussion. We thank the reviewer for pointing this out. Our original intention was to use the low-dimensional case as an illustrative starting point: in two dimensions, the pruning acceleration effect can be characterized geometrically, which is formalized in Lemma 3. In the high-dimensional setting, the same phenomenon is captured as *sublinear-time progress*, namely that a linear number of points can be processed within a sublinear number of iterations; this is formalized in Lemma 4.

We agree that separating these two results into different subsections makes the presentation clearer. In the revised manuscript, we have reorganized this part accordingly. Section 4.3 now contains only the two-dimensional base case in Lemma 3, while the new Section 4.4, “High-Dimensional Settings,” contains Lemma 4. We also added the following transition sentence at the beginning of Section 4.4:

We now extend the analysis to general d -dimensional space. The same phenomenon as in the 2D case manifests as *sublinear-time progress*: a linear number of points is processed within sublinear iterations. This behavior is formalized in the following lemma.

O2 / R2 – Clarification of the proof of Theorem 1

O2. The change from Lemma 4 (requiring uniform data distribution) to Theorem 1 (applicable to any data distribution) is significant. I could not understand the following parts in the proof of Theorem 1. In part (1), is it correct to use Lemma 4 (requiring uniform data distribution)? In part (2), please provide more details on how the Differential Equation Method is used properly to find the number of iterations.

R2. Clarify the proof of Theorem 1 (see O2).

What we changed: We rewrote the proof of Theorem 1 in 4.5 in three explicit cases and added a detailed verification of the **Differential Equation Method (DEM)** conditions. The revised proof first establishes the result under the uniform distribution model, where Lemma 4 applies, and then lifts the same argument to the general- α and non-uniform density settings.

Response and discussion. (1) We have substantially revised the proof of Theorem 1 to make the logical structure explicit. The key point is that **Lemma 4 is used as the technical foundation for the uniform-distribution case**, while Theorem 1 is obtained by identifying the exact condition required by the proof and then verifying that condition in more general settings.

(2) How the DEM is applied. We first clarify what the Differential Equation Method (DEM) refers to. Due to space limitations in the main text, the detailed statement of the DEM theorem is provided in the appendix.

Lemma 1 (Differential Equation Method). *Given integers $n \geq 1$, a bounded domain $\mathcal{D} \subseteq \mathbb{R}^{d+1}$, a function $F : \mathcal{D} \rightarrow \mathbb{R}$, and a sequence of sigma-algebras $\mathcal{F}_0 \subseteq \mathcal{F}_1 \subseteq \dots$, define random variables $Y(i)$ such that $Y(i)$ is $\mathcal{F}_i$ -measurable for $i \geq 0$.*

Assume that for all $i \geq 0$, whenever $(i/n, Y(i)/n) \in \mathcal{D}$, the following conditions are satisfied. In addition, assume that $\mathcal{D}$ contains the closure of the set $\{(0, z_0) : \mathbb{P}(Y(0) = z_0 n) \neq 0\}$, meaning that $\mathcal{D}$ includes all possible starting points.

If $Y(t)$ is a submartingale and the following three conditions hold:

(1) (Boundedness) *For some $\beta = \beta(n) \geq 1$ and $\gamma = \gamma(n)$, the probability that $|Y(i+1) - Y(i)| \leq \beta$ given $\mathcal{F}_i$ is at least $1 - \gamma$. This ensures that the one-step change of $Y(i)$ is controlled.*

(2) (Trend) *For some $\lambda_1 = \lambda_1(n) = o(1)$, the difference between the conditional expectation of the one-step change and the normalized value of the function F is negligible:*

$$\left| \mathbb{E}(Y(i+1) - Y(i) \mid \mathcal{F}_i) - F\left(\frac{i}{n}, \frac{Y(i)}{n}\right) \right| \leq \lambda_1.$$

(3) (Lipschitz) *F is continuous and Lipschitz with constant L on $\mathcal{D}$.*

Then we have:

(a) *For $(0, \hat{z}_0) \in \mathcal{D}$, the differential equation*

$$\frac{dz}{dx} = F(x, z)$$

has a unique solution in $\mathcal{D}$ with the initial condition $z(0) = \hat{z}_0$. This solution can be extended close to the boundary of $\mathcal{D}$.

Let $\lambda > \lambda_1 + C_0 n \gamma$, where $\lambda = o(1)$. Then, with probability

$$1 - O\left(n\gamma + \frac{\beta}{\lambda} \exp\left(-\frac{n\lambda^3}{\beta^3}\right)\right),$$

the following holds:

$$Y(i) = nz\left(\frac{i}{n}\right) + O(\lambda n),$$

where $z(x)$ is the solution from part (a) with initial condition $\hat{z}_0 = \frac{1}{n}Y(0)$. This result holds uniformly for $0 \leq i \leq \sigma n$, where $\sigma = \sigma(n)$ is the supremum of x such that the solution $z(x)$ remains within $\mathcal{D}$.

We also expand the explanation of the DEM part and make its role in the proof clearer. After the fast-pruning phase, the remaining process evolves with relatively small normalized one-step changes. We introduce a **normalized pruning process** Z_t and approximate its conditional drift by a deterministic function

$$F(z) = 1 + c_*(1 - z).$$

The revised proof explicitly checks the three standard Wormald DEM conditions:

• *Boundedness:*

$$|Y_{t+1} - Y_t| = 1 + \Delta S_{t+1} \leq |S_t| \leq n.$$

• *Trend:*

$$\mathbb{E}[Z_{t+1} - Z_t \mid \mathcal{F}_t] = \frac{F(Z_t) + o(1)}{n}, \quad F(z) = 1 + c_*(1 - z).$$

• *Lipschitz continuity:*

$$|F(z_1) - F(z_2)| = c_*|z_1 - z_2|.$$

These conditions justify approximating the discrete pruning dynamics by the corresponding differential equation in the plateau phase. The solution of this differential equation then gives the required bound on the remaining number of iterations. Combining the fast-pruning phase with the DEM-controlled plateau phase yields the final iteration bound

$$T = O(n^{2/3+\epsilon})$$

with probability 1. Since the out-degree is controlled by the number of pruning iterations, this proves the desired degree bound in Theorem 1.

Additional discussion on the proof details.

As established in Lemma 1.3, the out-degree of the constructed SNG can be controlled by the number of pruning iterations. Therefore, proving the degree bound reduces to proving an upper bound on the number of iterations required to prune all candidates. In particular, the goal is to show that the total number of iterations is at most

$$O(n^{2/3+\epsilon})$$

with probability 1, for any fixed small $\epsilon > 0$.

The proof is motivated by the two-phase pruning behavior observed in Section 4.1: an initial fast-pruning phase followed by a slower plateau phase. Lemma 4 formalizes this behavior in the uniform-distribution setting. Specifically, it first shows that within $n^{1/2}$ iterations, a higher-order number of points can be pruned; it then strengthens this statement in Lemma 4(2), showing that any sublinear number of iterations can remove nearly linearly many points. This result captures the fast-pruning behavior. After this phase, only a small residual set of candidates remains, and the one-step changes become sufficiently regular for the DEM approximation to apply.

Lemma 4 is indeed proved under the uniform-distribution assumption. In the revised proof, we use it as the base case and then make explicit which part of the argument is actually needed for the general theorem. The two-phase pruning analysis and DEM require a uniform positive lower bound on the per-iteration pruning probability. Therefore, the key task is to construct a constant

$$\pi_t \geq c_* > 0,$$

or, equivalently, to show that the pruning probability does not tend to zero. Once such a constant c_* is available, the Chernoff bound, the Borel–Cantelli argument, and the DEM step all proceed in the same way.

For completeness, we include the revised proof sketch of Theorem 1 below. The full proof, including the detailed calculations, is provided in Appendix B.3.1 of the full version available on our GitHub repository[1] and arxiv[2].

PROOF SKETCH. The proof has two ingredients. The first is a uniform positive lower bound on the per-iteration pruning probability, namely $\pi_t \geq c_* > 0$. The second is the two-phase pruning analysis developed in Sections 4.1–4.4, which combines the Chernoff bound, the Borel–Cantelli argument, and Wormald’s differential equation method (DEM). Since this analysis only requires the existence of some positive lower bound on π_t , it is insensitive to the specific distributional model and to the value of α . We therefore first prove the bound in the uniform $\alpha = 1$ setting, and then extend the same argument to general $\alpha > 0$ and to bounded non-uniform densities.

Case 1: uniform distribution and $\alpha = 1$. We first consider the base setting where the candidates are uniformly distributed and $\alpha = 1$. In this case, the blocking condition

$$\|p - x\| \geq \|p_t^* - x\|$$

defines a spherical-cap region. By Lemma 2, its normalized volume is bounded below by a positive constant, so

$$\pi_t \geq c(d, 1) > 0.$$

Let $Y_t := |S'_t|$ be the number of processed points after t iterations. Since $|S_t| = n - 1 - Y_t$, the one-step increment satisfies

$$Y_{t+1} - Y_t = 1 + \Delta S_{t+1},$$

where the term 1 corresponds to the selected nearest candidate and ΔS_{t+1} is the number of additionally pruned candidates. Therefore, conditional on the current filtration $\mathcal{F}_t$,

$$\mathbb{E}[Y_{t+1} - Y_t \mid \mathcal{F}_t] = 1 + \mathbb{E}[\Delta S_{t+1} \mid \mathcal{F}_t] \geq 1 + c(d, 1)(n - 1 - Y_t).$$

After normalization $Z_t := Y_t/n$, this becomes

$$\mathbb{E}[Z_{t+1} - Z_t \mid \mathcal{F}_t] \geq \frac{1 + c(d, 1)(n - 1 - nZ_t)}{n}.$$

Thus the associated deterministic trajectory is governed, up to lower-order terms, by the ODE

$$\frac{dz}{d\tau} = 1 + c(d, 1)(1 - z),$$

where $\tau = t/n$ is the rescaled time variable.

We verify the DEM conditions as follows. The boundedness condition follows from

$$|Y_{t+1} - Y_t| = 1 + \Delta S_{t+1} \leq |S_t| \leq n.$$

The trend condition follows from the conditional drift estimate above:

$$\mathbb{E}[Z_{t+1} - Z_t \mid \mathcal{F}_t] = \frac{F(Z_t) + o(1)}{n}, \quad F(z) := 1 + c(d, 1)(1 - z).$$

Finally, F is Lipschitz since

$$|F(z_1) - F(z_2)| = c(d, 1)|z_1 - z_2|.$$

Therefore Wormald’s DEM applies, and the stochastic trajectory of Z_t follows the solution of the above ODE with high probability.

Combining this DEM approximation with the Chernoff and Borel–Cantelli steps, we obtain the two-phase pruning bound. For any $\nu \in (2/3, 1)$, the process reaches

$$|S'_t| \geq n - n^{1-\nu}$$

within $t_1 = O(n^\nu)$ iterations with probability one. Once this level is reached, at most $n^{1-\nu}$ candidates remain. Applying the same DEM argument to the residual process, the plateau phase requires $O(n^{1-\nu})$ additional iterations to reach $|S'_t| = n - 1$. Hence

$$T = O(n^\nu) + O(n^{1-\nu}) = O(n^\nu).$$

Taking $\nu = 2/3 + \varepsilon$ gives

$$T = O(n^{2/3+\varepsilon}).$$

Case 2: uniform distribution and general $\alpha > 0$. For a general relaxation parameter α , the blocking condition becomes

$$\|p - x\| \geq \alpha \|p_t^* - x\|,$$

which defines an Apollonius-type blocking region. Although the geometry differs from the spherical-cap case, its normalized volume inside the current local support is still bounded below by a positive constant. Thus there exists

$$c(d, \alpha) > 0$$

such that

$$\pi_t \geq c(d, \alpha).$$

The DEM proof above uses the pruning geometry only through the positive constant in the drift lower bound. Hence the same argument applies after replacing $c(d, 1)$ by $c(d, \alpha)$. The drift function becomes

$$F_\alpha(z) := 1 + c(d, \alpha)(1 - z),$$

which is still Lipschitz, while the boundedness and trend conditions remain unchanged. Therefore,

$$T = O(n^{2/3+\varepsilon})$$

also holds for every fixed $\alpha > 0$.

Case 3: bounded non-uniform density. Suppose the local candidate distribution has density f satisfying $0 < m \leq f \leq M < \infty$. This condition captures many locally regular non-uniform settings while excluding degenerate cases with vanishing density. Conditional on p_t^* , the blocking region $R_t(p_t^*)$ from Case 2 gives

$$\pi_t \geq \frac{m}{M} \frac{\text{Vol}(R_t(p_t^*))}{\text{Vol}(\Omega_{t-1})} \geq \frac{m}{M} c(d, \alpha) =: c_* > 0.$$

Thus non-uniformity only changes the constant in the drift lower bound. The DEM drift becomes

$$F_*(z) := 1 + c_*(1 - z),$$

which remains Lipschitz, and the boundedness and trend conditions are unchanged. Hence the Chernoff, Borel–Cantelli, and DEM steps still yield

$$T = O(n^{2/3+\varepsilon}).$$

□

O3 / R3 – Applicability of Theorem 2 to general distributions

O3. Theorem 2 is designed for any given data distribution. However, in the proof of Theorem 2, why is there an assumption that “they are uniformly distributed in a d -dimensional ball”?

R3. Is Theorem 2 applicable to any given data distribution OR the uniform distribution only?

What we changed: We now state Theorem 2 explicitly under the uniform-distribution assumption and discuss its extension to non-uniform distributions as future work. [Sec. 5 Theorem 2 statement; Sec. 9, newly added third paragraph in the Conclusions]

Response and discussion. We thank the reviewer for pointing this out. We agree that the previous wording could be read as suggesting applicability to arbitrary data distributions, which was not sufficiently precise. We have therefore revised the statement and discussion of Theorem 2 to clarify its scope. In the current version, Theorem 2 is rigorously proved under the uniform-distribution assumption.

We believe that extending the result beyond the uniform case is a promising direction. Our DiskANN experiments show that the empirical $O(\log n)$ search-path behavior persists across real datasets, which are clearly non-uniform. However, a rigorous proof for non-uniform distributions would require establishing the corresponding probability inequalities and drift conditions under the induced distribution. We therefore identify this extension as future work in the discussion section.

O4 – Recall improvement under OPT

O4. It is unclear why the proposed parameter optimization in Section 6 can slightly improve the recall in the experiment of Table 4. I could not find any formulation of the recall in Section 6.

What we changed: We expanded the search-performance discussion to explain why the analytically selected truncation parameter R^* can slightly improve Recall@10. In particular, we clarify how both overly small and overly large truncation parameters can hurt search performance. [Sec. 7.3, third paragraph, search-performance discussion]

Response and discussion. We appreciate the reviewer’s careful observation. In the original manuscript, we briefly explained that OPT consistently matches or improves Recall@10 across all three algorithms and all datasets, with the improvements being most pronounced on DEEP1M, MSong, GIST1M, and GloVe. We attributed this effect to the fact that analytical parameter selection can avoid locally optimal but globally suboptimal truncation choices. We agree that this explanation was too brief, and we have now expanded the discussion.

The mechanism is as follows. The truncation parameter R controls the trade-off between graph connectivity and construction/search overhead.

When $R \ll R^*$, the truncation is too aggressive. Many effective SNG edges are artificially removed, which makes the graph under-connected. As a result, greedy search requires more hops and is




Figure 1: Query performance on billion-scale datasets

more likely to miss the true nearest neighbors, leading to lower recall.

When $R \gg R^*$, truncation is no longer the limiting factor, but the construction procedure pays additional overhead for maintaining R -NN candidates and processing reverse edges. These redundant edges slow down each search step, but they do not necessarily provide additional useful connectivity, so recall no longer improves and may even degrade under a fixed latency or search-budget setting.

Therefore, both overly small and overly large choices of R can hurt search performance. The analytically derived R^* provides a better balance between preserving useful SNG edges and avoiding redundant construction/search overhead. In contrast, a discrete parameter sweep may terminate at a locally optimal but globally suboptimal truncation value, especially when the recall-latency surface is non-convex or noisy. We believe the slight recall improvement observed in Table 4 comes from this calibration effect: the proposed R^* helps correct suboptimal sweep choices and places the graph closer to the effective operating regime determined by the structural parameter α .

We have added this explanation to the search-performance discussion in the revised manuscript.

O5 – Scalability on billion-scale datasets

O5. In the experimental study, the dataset size is up to 1.2 million only. However, in the TriBase paper (ACM on Management of Data 2025), several datasets with 1 billion vectors (e.g., DEEP1B, SpaceV1B, SIFT1B) are used in the experimental study. In order to test the scalability, the authors are recommended to include these huge datasets in the experimental study.

What we changed: We added a billion-scale evaluation on three standard benchmarks, SIFT1B, DEEP1B, and SpaceV1B, following the protocol of recent large-scale ANNS evaluations such as TriBase. The results are reported in a new figure. [New Sec. 7.6, billion-scale block, third–fourth paragraphs; Figure 5]

Response and discussion. We thank the reviewer for this suggestion. Adding billion-scale datasets makes the empirical evaluation more complete and provides a more direct test of scalability beyond the million-scale setting. Accordingly, we added a new evaluation block on three standard billion-scale benchmarks: SIFT1B, DEEP1B, and SpaceV1B.

Across all three datasets, OPT-SNG preserves the construction-time advantage observed at the million-point scale. In particular, OPT-Vamana achieves graph-construction speedups of 2.48 $\times$ on SIFT1B, 1.03 $\times$ on DEEP1B, and 1.07 $\times$ on SpaceV1B. Moreover, at fixed latency targets, OPT-SNG matches or improves Recall@10 compared with the original version on all three datasets.




Figure 2: Given that p^* is the nearest point to p in t -th iteration and $\|p - p^*\| = \rho_t$. The light grey region represents all p' satisfied pruning condition $\|p - p'\| > \|p^* - p'\|$.

Due to hardware constraints, the absolute QPS in our billion-scale experiments is lower than what may be achieved on large-scale production systems. However, all methods are evaluated under the same hardware setting, and the comparison therefore still reflects the relative effectiveness of the proposed analytical parameter selection. These results demonstrate that OPT-SNG remains effective at billion-vector scale.

D1 – Missing proof of Lemma 2

D1. The proof of Lemma 2 is missing.

What we changed and response: In the original manuscript, the second paragraph of Section 4 stated that, for completeness, all detailed proofs of the lemmas and theorems are provided in the full version available on our GitHub repository [1] and arXiv [2]. We thank the reviewer for pointing out that the location of the proof of Lemma 2 was not sufficiently clear. Since Lemma 2 is a supporting lemma rather than a main theorem, its full proof is not included in the main text due to space limitations. To avoid ambiguity, we have added an explicit cross-reference immediately after the statement of Lemma 2, directing readers to the appendix of the arXiv/full version[2]. [Sec. 4.2, fourth paragraph, Lemma 2 statement]

For completeness, we also reproduce the proof below.

Lemma 2 (Pruning Probability in SNG Construction). *Given a dataset $P \subset \mathbb{R}^d$ consisting of n points independently and uniformly distributed within a ball of radius ρ_0 , centered at p , we consider the SNG construction processed on the center point p with $\alpha = 1$. The probability π_t that a point p' is pruned in the t -th iteration is given by*

$$\pi_t = \frac{I_{1-(\rho_t/2\rho_0)^2} \left(\frac{d+1}{2}, \frac{1}{2} \right) - (\rho_t/\rho_0)^d \cdot I_{0.75} \left(\frac{d+1}{2}, \frac{1}{2} \right)}{2(1 - (\rho_t/\rho_0)^d)},$$

where $I_x(a, b)$ denotes the regularized incomplete Beta function, and ρ_t is the distance between the nearest neighbor and p in t -th iteration. For large d , this probability satisfies

$$\frac{I_{0.75} \left(\frac{d+1}{2}, \frac{1}{2} \right)}{2} < \pi_t < \frac{1}{2}.$$

PROOF OF LEMMA 2. Assuming a uniform distribution over the search space, the probability that a randomly sampled point falls within a region is proportional to the region's volume. Without loss of generality, we place the center point p at the origin and align

the vector $\overrightarrow{pp^*}$ with the last coordinate axis, as shown in Figure 2. Under this coordinate system, the nearest neighbor p^* lies at the position $(0, 0, \dots, 0, \rho_t) \in \mathbb{R}^d$. Let p^* denote the nearest point to p in t -th iteration and $\|p - p^*\| = \rho_t$, then the hyperplane is defined by:

$$\omega \cdot x + b = 0, \quad \text{where } \omega = (0, \dots, 0, 1), \quad b = -\frac{\rho_t}{2}.$$

The pruned region at iteration t , denoted by $D(t)$, is defined as the intersection of the ball $B(p, \rho_0)$ and the half-space $\omega \cdot x > \frac{\rho_t}{2}$. Let $V(\cdot)$ denote the Lebesgue volume in the appropriate dimension. For example, $V(B_d(r))$ is the volume of a d -dimensional ball of radius r , and $V(B_{d-1}(r))$ is the volume of a $(d-1)$ -dimensional ball of radius r . In particular, $V(B_d(1))$ denotes the volume of the d -dimensional unit ball:

$$V(B_d(1)) = \frac{\pi^{d/2}}{\Gamma(1 + d/2)}.$$

To compute $V(D(t))$, we slice the ball orthogonally along the last coordinate x_d , and integrate over the $(d-1)$ -dimensional cross-sectional volumes. The volume of the pruned region is given by:

$$V(D(t)) = \int_{x_d=\frac{\rho_t}{2}}^{\rho_0} V\left(B_{d-1}\left(\sqrt{\rho_0^2 - x_d^2}\right)\right) dx_d,$$

where $B_{d-1}(\sqrt{\rho_0^2 - x_d^2})$ is the $(d-1)$ -dimensional ball of radius $\sqrt{\rho_0^2 - x_d^2}$, corresponding to the horizontal slice at height x_d . Using the change of variables $t = x_d/\rho_0$, we have $x_d = \rho_0 t$ and $dx_d = \rho_0 dt$. Substituting this into the integral yields:

$$V(D(t)) = V(B_{d-1}(1)) \cdot \rho_0^d \int_{\frac{\rho_t}{2\rho_0}}^1 (1 - t^2)^{\frac{d-1}{2}} dt.$$

Next, we apply the substitution $z = 1 - t^2$ to convert the integral into a standard form. Under this change of variables, we obtain:

$$\int_{\frac{\rho_t}{2\rho_0}}^1 (1 - t^2)^{\frac{d-1}{2}} dt = \frac{1}{2} \int_0^{1-(\frac{\rho_t}{2\rho_0})^2} z^{\frac{d-1}{2}} (1 - z)^{-1/2} dz.$$

This integral corresponds to the incomplete Beta function. After normalization, it becomes proportional to the *regularized incomplete Beta function*, also known as the cumulative distribution function of the Beta distribution, denoted by $I_x(a, b)$. Thus, we have:

$$\int_{\frac{\rho_t}{2\rho_0}}^1 (1 - t^2)^{\frac{d-1}{2}} dt = \frac{\Gamma\left(\frac{d+1}{2}\right) \Gamma\left(\frac{1}{2}\right)}{2\Gamma\left(\frac{d}{2} + 1\right)} \cdot I_{1-(\frac{\rho_t}{2\rho_0})^2} \left(\frac{d+1}{2}, \frac{1}{2} \right),$$

where $\Gamma(\cdot)$ denotes the Gamma function. Hence, the volume of the pruned region is:

$$V(D(t)) = \frac{\Gamma\left(\frac{d+1}{2}\right) \Gamma\left(\frac{1}{2}\right)}{2\Gamma\left(\frac{d}{2} + 1\right)} V(B_{d-1}(1)) \cdot \rho_0^d \cdot I_{1-(\frac{\rho_t}{2\rho_0})^2} \left(\frac{d+1}{2}, \frac{1}{2} \right).$$

However, since this includes points within the inner ball $B(\rho_t)$, and we subtract the volume within $B(\rho_t)$, denoted V_{inner} , is:

$$V_{\text{inner}} = \frac{\Gamma\left(\frac{d+1}{2}\right) \Gamma\left(\frac{1}{2}\right)}{2\Gamma\left(\frac{d}{2} + 1\right)} V(B_{d-1}(1)) \cdot \rho_t^d \cdot I_{3/4} \left(\frac{d+1}{2}, \frac{1}{2} \right),$$

where $3/4 = 1 - (1/2)^2$ accounts for the halfway-positioned hyperplane. The probability that a randomly sampled point is pruned in iteration t is:

$$\pi_t = \frac{V(D(t)) - V_{\text{inner}}}{V(B_d(\rho_0)) - V(B_d(\rho_t))}.$$

Substituting the expressions above and simplifying using $V(B_d(1)) = \frac{\pi^{d/2}}{\Gamma(\frac{d}{2}+1)}$, we get:

$$\pi_t = \frac{I_{1-(\frac{\rho_t}{\rho_0})^2} - (\frac{\rho_t}{\rho_0})^d \cdot I_{3/4}}{2 \left(1 - (\frac{\rho_t}{\rho_0})^d\right)}.$$

Since $\frac{\rho_t}{\rho_0} \in (0, 1)$ and $1 - (\frac{\rho_t}{\rho_0})^2 \in (3/4, 1)$, the ratio is always decreases with respect to $\frac{\rho_t}{\rho_0}$. Specifically, noting that $I_{3/4}(a, b) < 1$, for large d , we derive:

$$\frac{1}{2} I_{3/4} \left(\frac{d+1}{2}, \frac{1}{2} \right) < \pi_t < \frac{1}{2}.$$

2-dimensional case. In the two-dimensional case, we can leverage geometric relationships to derive a lower bound of $\frac{1}{3}$ for the pruning probability. This bound corresponds to the scenario where the pruning probability decreases as ρ_t/ρ_0 increases, reaching its minimum when $\rho_t/\rho_0 = 1$ —that is, when the distance from the center point p to its nearest neighbor p^* approaches the radius of the disk. In this situation, the pruning decision boundary can be approximated by the perpendicular bisector of the segment connecting p and p^* . Geometrically, this bisector divides the circular region into two sectors, and the prunable region corresponds to the sector opposite p . As the angle subtended by this region exceeds $\frac{2\pi}{3}$, it follows that the fraction of the total area being pruned is at least $\frac{1}{3}$. Thus, the pruning probability is bounded below by $\frac{1}{3}$ in the two-dimensional setting. $\square$

D2 — Notation and “almost surely” in the proof of Lemma 4

D2. In the proof of Lemma 4, what is ?? Please clarify the meaning of “almost surely” in the proof.

What we changed:

- We corrected the missing-reference marker “??” in the proof of Lemma 4. [Sec. 4.4, proof of Lemma 4, first line below Eq. (12)]
- We added a footnote at the first occurrence of “almost surely” to clarify its measure-theoretic meaning. [Sec. 4.4, proof of Lemma 4, third line]

Response and discussion. We thank the reviewer for pointing out these issues. The “??” marker was caused by an unresolved cross-reference in the compiled manuscript. It was intended to refer to Lemma 5 in Appendix A, and this reference has now been corrected. We also agree that the phrase “almost surely” may be unfamiliar to some readers outside probability theory. Since this is the standard terminology used in measure-theoretic probability, we keep the phrase but add the following explanatory footnote at its first occurrence:

“Throughout this paper, ‘almost surely’ (abbreviated *a.s.*) is used in its standard measure-theoretic sense: the event

holds with probability 1, i.e., the set of outcomes on which it fails has probability measure zero.”

RESPONSE TO REVIEWER #2

We thank the reviewer for the thoughtful and technically precise feedback. We respond to each concern in detail below.

What we changed:

- We rewrote the proof of Theorem 1 in three explicit cases and clarified how the result extends beyond the uniform $\alpha = 1$ base case. [Sec. 4.5, proof of Theorem 1]

Response and discussion. We thank the reviewer for pointing out this issue. In the revised manuscript, we make the extension argument explicit. The key observation is that the two-phase pruning analysis does not rely on the exact uniform-density formula itself; it only requires a uniform positive lower bound on the per-iteration pruning probability. We therefore organize the proof into the following three cases:

- **Uniform distribution with $\alpha = 1$:** Lemma 2 provides the base pruning-probability lower bound $\pi_t \geq c(d, 1) > 0$.
- **Uniform distribution with general $\alpha > 0$:** the spherical cap is replaced by the corresponding Apollonius blocking region, yielding a lower bound $\pi_t \geq c(d, \alpha) > 0$.
- **Bounded non-uniform density:** if the local density satisfies $0 < m \leq f \leq M < \infty$, the pruning probability remains bounded below by $\pi_t \geq (m/M)c(d, \alpha) =: c_* > 0$.

After this lower bound is established in each case, the same Chernoff, Borel–Cantelli, and DEM argument applies. The dimension d enters the analysis through the pruning-probability estimate and is absorbed into the constants $c(d, 1)$ and $c(d, \alpha)$, so the argument also applies to high-dimensional data.

For clustered or manifold-like data, the same proof strategy would require an appropriate lower bound on the local pruning probability under the induced data distribution. One possible route is to approximate the local distribution by a regular density model, such as a mixture model, and then verify the corresponding local pruning-probability lower bound. We include this discussion to clarify how the analysis may be adapted beyond the idealized uniform setting.

For completeness, we include the revised proof sketch of Theorem 1 below. The full proof, including the detailed calculations, is provided in Appendix B.3.1 of the full version available on our GitHub repository [1] and arXiv [2].

PROOF SKETCH. The proof has two ingredients. The first is a uniform positive lower bound on the per-iteration pruning probability, namely $\pi_t \geq c_* > 0$. The second is the two-phase pruning analysis developed in Sections 4.1–4.4, which combines the Chernoff bound, the Borel–Cantelli argument, and Wormald’s differential equation method (DEM). Since this analysis only requires the existence of some positive lower bound on π_t , it is insensitive to the specific distributional model and to the value of α . We therefore first prove the bound in the uniform $\alpha = 1$ setting, and then extend the same argument to general $\alpha > 0$ and to bounded non-uniform densities.

Case 1: uniform distribution and $\alpha = 1$. We first consider the base setting where the candidates are uniformly distributed and

$\alpha = 1$. In this case, the blocking condition

$$\|p - x\| \geq \|p_t^* - x\|$$

defines a spherical-cap region. By Lemma 2, its normalized volume is bounded below by a positive constant, so

$$\pi_t \geq c(d, 1) > 0.$$

Let $Y_t := |S'_t|$ be the number of processed points after t iterations. Since $|S_t| = n - 1 - Y_t$, the one-step increment satisfies

$$Y_{t+1} - Y_t = 1 + \Delta S_{t+1},$$

where the term 1 corresponds to the selected nearest candidate and ΔS_{t+1} is the number of additionally pruned candidates. Therefore, conditional on the current filtration $\mathcal{F}_t$,

$$\mathbb{E}[Y_{t+1} - Y_t \mid \mathcal{F}_t] = 1 + \mathbb{E}[\Delta S_{t+1} \mid \mathcal{F}_t] \geq 1 + c(d, 1)(n - 1 - Y_t).$$

After normalization $Z_t := Y_t/n$, this becomes

$$\mathbb{E}[Z_{t+1} - Z_t \mid \mathcal{F}_t] \geq \frac{1 + c(d, 1)(n - 1 - nZ_t)}{n}.$$

Thus the associated deterministic trajectory is governed, up to lower-order terms, by the ODE

$$\frac{dz}{d\tau} = 1 + c(d, 1)(1 - z),$$

where $\tau = t/n$ is the rescaled time variable.

We verify the DEM conditions as follows. The boundedness condition follows from

$$|Y_{t+1} - Y_t| = 1 + \Delta S_{t+1} \leq |S_t| \leq n.$$

The trend condition follows from the conditional drift estimate above:

$$\mathbb{E}[Z_{t+1} - Z_t \mid \mathcal{F}_t] = \frac{F(Z_t) + o(1)}{n}, \quad F(z) := 1 + c(d, 1)(1 - z).$$

Finally, F is Lipschitz since

$$|F(z_1) - F(z_2)| = c(d, 1)|z_1 - z_2|.$$

Therefore Wormald's DEM applies, and the stochastic trajectory of Z_t follows the solution of the above ODE with high probability.

Combining this DEM approximation with the Chernoff and Borel–Cantelli steps, we obtain the two-phase pruning bound. For any $\nu \in (2/3, 1)$, the process reaches

$$|S'_t| \geq n - n^{1-\nu}$$

within $t_1 = O(n^\nu)$ iterations with probability one. Once this level is reached, at most $n^{1-\nu}$ candidates remain. Applying the same DEM argument to the residual process, the plateau phase requires $O(n^{1-\nu})$ additional iterations to reach $|S'_t| = n - 1$. Hence

$$T = O(n^\nu) + O(n^{1-\nu}) = O(n^\nu).$$

Taking $\nu = 2/3 + \varepsilon$ gives

$$T = O(n^{2/3+\varepsilon}).$$

Case 2: uniform distribution and general $\alpha > 0$. For a general relaxation parameter α , the blocking condition becomes

$$\|p - x\| \geq \alpha \|p_t^* - x\|,$$

which defines an Apollonius-type blocking region. Although the geometry differs from the spherical-cap case, its normalized volume

inside the current local support is still bounded below by a positive constant. Thus there exists

$$c(d, \alpha) > 0$$

such that

$$\pi_t \geq c(d, \alpha).$$

The DEM proof above uses the pruning geometry only through the positive constant in the drift lower bound. Hence the same argument applies after replacing $c(d, 1)$ by $c(d, \alpha)$. The drift function becomes

$$F_\alpha(z) := 1 + c(d, \alpha)(1 - z),$$

which is still Lipschitz, while the boundedness and trend conditions remain unchanged. Therefore,

$$T = O(n^{2/3+\varepsilon})$$

also holds for every fixed $\alpha > 0$.

Case 3: bounded non-uniform density. Suppose the local candidate distribution has density f satisfying $0 < m \leq f \leq M < \infty$. This condition captures many locally regular non-uniform settings while excluding degenerate cases with vanishing density. Conditional on p_t^* , the blocking region $R_t(p_t^*)$ from Case 2 gives

$$\pi_t \geq \frac{m}{M} \frac{\text{Vol}(R_t(p_t^*))}{\text{Vol}(\Omega_{t-1})} \geq \frac{m}{M} c(d, \alpha) =: c_* > 0.$$

Thus non-uniformity only changes the constant in the drift lower bound. The DEM drift becomes

$$F_*(z) := 1 + c_*(1 - z),$$

which remains Lipschitz, and the boundedness and trend conditions are unchanged. Hence the Chernoff, Borel–Cantelli, and DEM steps still yield

$$T = O(n^{2/3+\varepsilon}).$$

□

D2 – Independence assumption underlying the Binomial model

D2. In the Lemma 3 proof sketch, the authors state “remaining points are pruned independently with probability π_i ” and then they model ΔS_i as binomial. However, in SNG pruning, geometric blocking events are typically highly dependent — in particular, they are correlated via distances and shared neighbors, especially in high dimension and non-uniform data.

What we changed: We revised the proof of Lemma 3 to clarify that the Binomial model is used through a *conditional independence* argument. In particular, the independence is taken conditional on the currently selected nearest neighbor p_t^* , rather than as an unconditional independence assumption over the whole pruning process. [Sec. 4.3, proof of Lemma 3, first paragraph, third line]

Response and discussion. We thank the reviewer for raising this important point. The independence used in Lemma 3 is more limited and should be understood conditionally with respect to the natural filtration of the pruning process. At a fixed iteration t , once




Figure 3: Ablation: Parameter reuse across datasets.

the current nearest neighbor p_t^* is fixed, the pruning event for a remaining candidate p' ,

$$A_{p'} = \{\|p - p'\| \geq \alpha \|p_t^* - p'\|\},$$

depends only on the relative position of this single candidate p' with respect to the two fixed reference points (p, p_t^*) . Under the sampling model used in Lemma 3, the remaining candidates are conditionally i.i.d.; hence their relative positions to (p, p_t^*) do not influence one another. Therefore, for this fixed iteration and conditional on p_t^* , the events $\{A_{p'}\}$ are mutually independent, and the number of additionally pruned points satisfies

$$\Delta S_t \mid p_t^* \sim \text{Bin}(|S_{t-1}| - 1, \pi_t).$$

We have added this clarification directly in the proof of Lemma 3. This revision makes explicit that the Binomial model is not intended to assert independence across different pruning iterations or across the entire graph-construction process. It is only used as a one-step conditional model within a fixed iteration, which is sufficient for the drift and concentration estimates used in the proof.

We thank the reviewer for raising this important concern. To address it, we added the requested ablation study.

Parameter-reuse setting. In the new ablation, the baseline first performs a parameter sweep over R on one dataset and then reuses the selected value on subsequent builds from the same dataset family. Thus, the reuse baseline pays no additional tuning cost on the new builds. This setting is favorable to the baseline and directly tests whether OPT-SNG still provides construction-time benefits after excluding repeated sweep overhead.

Result. As shown in Figure 3, OPT-SNG still consistently outperforms the reuse baselines across all five datasets (SIFT1M, GIST1M, DEEP1M, MSong, and GloVe) and all three index types (DiskANN, NSG, and HNSW). The average construction-time reductions are:

- DiskANN-Vamana: 1.09×
- NSG: 1.65×
- HNSW: 1.74×

Although the speedup becomes smaller after amortizing the tuning cost, OPT-SNG still provides direct construction-time improvements. More importantly, this ablation confirms that the reported gains are not merely an artifact of charging parameter-sweep overhead to the baseline. Instead, they also stem from the improved

quality of the analytically selected R^* , which remains beneficial even under parameter reuse.

RESPONSE TO REVIEWER #3

We thank the reviewer for the balanced assessment and the actionable suggestions. We address the reviewer’s comments point by point below.

W1 / D1 — Limited baselines

W1. While the theoretical derivation is robust, the empirical evaluation is somewhat narrow in scope regarding the competitive landscape.

D1. Limited baselines: The experiments are restricted to optimizing three specific algorithms: Vamana (DiskANN), NSG, and HNSW. While these are popular, the evaluation misses several other state-of-the-art competitors that currently dominate various ANNS benchmarks, e.g., VSAG, SymphonyQG, and others.

What we changed:

- We added four additional baselines covering major families of recent ANNS methods: **VSAG**, **SymphonyQG**, **RaBitQ**, and **DET-LSH**. [Sec. 7.1.6 “Additional baselines”]
- We added Recall-QPS and construction-time comparisons against these baselines. [Sec. 7.3, fifth new paragraph “Comparison with Additional Baselines”]
- We added the corresponding bibliography entries.

Response and discussion. We thank the reviewer for this helpful suggestion. We agree that broadening the experimental comparison makes the evaluation more complete and better situates our method within the current ANNS landscape. In the revised manuscript, we add four further baselines, covering graph-based methods (VSAG and SymphonyQG), a quantization-based method (RaBitQ), and a hashing-based method (DET-LSH). The added comparison is shown in Figure 4.

Among the added baselines, VSAG and SymphonyQG are the most competitive, especially at high recall on SIFT1M and Deep1M. Nevertheless, the OPT variants match or surpass their Recall-QPS trade-offs across the evaluated datasets, while requiring substantially less index-construction time. On average, the OPT variants build the index about 1.2× faster than VSAG and 1.8× faster than SymphonyQG. The non-graph baselines show different behavior: RaBitQ and DET-LSH degrade more noticeably in the high-recall regime, which is typically the most relevant setting for practical ANNS workloads. Overall, these additional comparisons place OPT-SNG more clearly within the current ANNS landscape and show that its advantage is not limited to the original three graph baselines.

W2 / D2 — Model assumptions and clustered data

W2. The assumptions made in the martingale model may not hold in all real-world scenarios, potentially limiting the generalizability of the results.




Figure 4: Comparison with additional baselines on six datasets.

D2. Model assumptions: Real-world high-dimensional data (like text embeddings or image vectors) is almost never uniformly distributed. It typically lies on lower-dimensional manifolds and exhibits strong clustering (“clumpiness”). The derived bounds ($O(n^{2/3+\epsilon})$ degree and $O(\log n)$ search) may not hold for clustered data. In highly clustered datasets, the “pruning” efficiency often drops because neighbors are geometrically closer and less likely to “block” each other effectively compared to a uniform distribution. The authors are encouraged to discuss how their model might adapt or be extended to account for such data distributions.

What we changed:

- We reorganized the proof of the degree bound in Theorem 1 into three cases and clarified how the argument extends from the uniform setting to bounded non-uniform densities. [Sec. 4.5; revised proof structure and discussion]
- We added a discussion clarifying the scope of Theorem 2: the current logarithmic search-path bound is rigorously proved under the uniform-distribution assumption, while the non-uniform extension is left as future work. [Sec. 9, third paragraph]

Response and discussion. We thank the reviewer for raising this important concern. In the revised manuscript, we clarify the scope of our theoretical results and distinguish the degree bound in Theorem 1 from the search-path bound in Theorem 2.

For the degree bound in Theorem 1, the proof starts from the uniform distribution as the geometric base case, but the key quantity needed by the subsequent concentration and DEM arguments is a positive lower bound on the per-iteration pruning probability. Once such a bound is available, the remaining steps follow the same structure as in the uniform case. We therefore reorganized the proof into three cases: the uniform case with $\alpha = 1$, the uniform case with general $\alpha > 0$, and the bounded non-uniform density case.

For clustered or manifold-like data, the same proof strategy would require an appropriate lower bound on the local pruning probability under the induced data distribution. One possible route is to approximate the local distribution by a regular density model, such as a mixture model, and then verify the corresponding pruning-probability lower bound.

For Theorem 2, we revised the statement and discussion to make its scope explicit. The current $O(\log n)$ search-path bound is rigorously proved under the uniform-distribution assumption. At the same time, extending Theorem 2 beyond the uniform case remains a promising direction. Our DiskANN experiments suggest that the empirical $O(\log n)$ search-path behavior persists on real non-uniform datasets. However, **a rigorous proof would require establishing the corresponding probability inequalities and**




Figure 5: Query performance on billion-scale datasets.

drift conditions under the induced non-uniform distribution. Since this is technically separate from the current. We identify it as future work in the discussion section.

W3 – Billion-scale evaluation

W3. Given that many real-world applications involve datasets with billions of vectors, the lack of evaluation at this scale raises questions about the framework’s performance and scalability in truly large-scale settings.

What we changed: We added a billion-scale evaluation on three standard benchmarks, SIFT1B, DEEP1B, and SpaceV1B, following the protocol of recent large-scale ANNS evaluations such as TriBase. The results are reported in a new figure. [New Sec. 7.6, billion-scale block, third–fourth paragraphs; Figure 5]

Response and discussion. We thank the reviewer for this suggestion. Adding billion-scale datasets makes the empirical evaluation more complete and provides a more direct test of scalability beyond the million-scale setting. Accordingly, we added a new evaluation block on three standard billion-scale benchmarks: SIFT1B, DEEP1B, and SpaceV1B.

Across all three datasets, OPT-SNG preserves the construction-time advantage observed at the million-point scale. In particular, OPT-Vamana achieves graph-construction speedups of 2.48× on SIFT1B, 1.03× on DEEP1B, and 1.07× on SpaceV1B. Moreover, at fixed latency targets, OPT-SNG matches or improves Recall@10 compared with the original version on all three datasets.

Due to hardware constraints, the absolute QPS in our billion-scale experiments is lower than what may be achieved on large-scale production systems. However, all methods are evaluated under the same hardware setting, and the comparison therefore still reflects the relative effectiveness of the proposed analytical parameter selection. These results demonstrate that OPT-SNG remains effective at billion-vector scale.

D3 – Paper organization: prioritize Theorem 1 and elaborate R^*

D3. Regarding the organization of the paper, I recommend prioritizing the full explanation of Theorem 1 in the main text over the proof and discussion of Lemma 3. Additionally, the derivation of R^ would benefit from further elaboration, as the cost model is currently a bit complex to follow.*

What we changed:

- We added a “Preview of the main result” paragraph at the beginning of Section 4, so that Theorem 1 is stated before the supporting lemmas. [Sec. 4, second paragraph]
- We rewrote Section 6.2 to derive R^* step by step from the construction-cost model. [Sec. 6.2, second paragraph]
- We added an explanation of why the analytically derived R^* can also slightly improve Recall by avoiding locally suboptimal truncation choices. [Sec. 7.3, third paragraph, search-performance discussion]

Response and discussion. We thank the reviewer for these helpful suggestions. We made three revisions to improve the organization and readability of the manuscript.

Theorem 1 is now prioritized in the main text. At the beginning of Section 4, we added the following preview paragraph so that readers can see the main conclusion before the technical lemmas:

A preview of the main result. Before delving into the technical analysis, we state the central conclusion of this section: *with probability 1, the maximum out-degree of the constructed SNG is $O(n^{2/3+\varepsilon})$ for any small $\varepsilon > 0$* (Theorem 1, Section 4.5). The remainder of Section 4 formalizes the two-phase pruning dynamics needed for this bound: Lemma 2 (pruning probability), Lemma 3 (2D base case), and Lemma 4 (sublinear progress in d dimensions).

This revision makes clear that Lemma 3 serves as the 2D geometric base case, while Theorem 1 is the main result of the section.

The derivation of R^ is now expanded.* We rewrote Section 6.2 to separate the modeling, optimization, and calibration steps. The construction cost is summarized as

$$n(C_1 R \log n + C_2 R^2 \log n / \alpha + C_3 \alpha R^3 + O(1)),$$

where the three terms correspond to initial candidate collection, SNG pruning, and reverse-edge handling, respectively. Treating R as a function of α , the dominant stationary condition is

$$-\frac{C_2 R^2 \log n}{\alpha^2} + C_3 R^3 = 0,$$

which gives

$$R^* \propto \frac{\log n}{\alpha^2}.$$

The hidden constant absorbs the dimension, implementation constants, and lower-order terms. Our algorithm then calibrates this constant using a single reference build at $R = n^{2/3}$ under α_1 , measures the average out-degree $\bar{R}$, sets

$$K' = \frac{\alpha_1^2 \bar{R}}{\log n},$$

and returns

$$R^* = K' \frac{\log n}{\alpha_2^2}$$

for the target α_2 . This avoids costly parameter sweeping while using the theoretical degree bound in Theorem 1 as a near-untruncated reference.

We also clarified the effect of R^ on Recall.* We added a short explanation of why this calibration may slightly improve Recall. When $R \ll R^*$, truncation is too aggressive and removes useful SNG edges, reducing connectivity and Recall. When $R \gg R^*$, maintaining R -NN candidates and reverse edges mainly introduces redundant edges and extra overhead, slowing search without further Recall gains. Thus, the analytically derived R^* can slightly improve Recall by avoiding locally suboptimal choices in discrete parameter sweeps.

We are grateful for the reviewers’ time and thoughtful engagement, and we look forward to the committee’s further evaluation.

REFERENCES

- [1] [n.d.]. OPT-SNG: Sparse Neighborhood Graph-Based Approximate Nearest Neighbor Search Revisited. <https://github.com/maxila320/OPT-SNG>. GitHub repository.
- [2] Xinran Ma, Zhaoqi Zhou, Chuan Zhou, Qi Meng, Zaijiu Shang, Guoliang Li, and Zhiming Ma. 2025. Graph-Based Approximate Nearest Neighbor Search Revisited: Theoretical Analysis and Optimization. *arXiv preprint arXiv:2509.15531* (2025). arXiv:2509.15531 [cs.DS]

Sparse Neighborhood Graph-Based Approximate Nearest Neighbor Search Revisited: Theoretical Analysis and Optimization

Xinran Ma*

Academy of Mathematics and
Systems Science, Chinese Academy of
Sciences
maxinran22@mails.ucas.ac.cn

Zhaoqi Zhou*

Huawei Technologies Co., Ltd.
zhouzhaoqi1@huawei.com

Chuan Zhou

Academy of Mathematics and
Systems Science, Chinese Academy of
Sciences
zhouchuan@amss.ac.cn

Zaijiu Shang

Shanghai Institute for Mathematics
and Interdisciplinary Sciences
zaijiu@simis.cn

Guoliang Li

Tsinghua University
liguoliang@tsinghua.edu.cn

Zhiming Ma

Academy of Mathematics and
Systems Science, Chinese Academy of
Sciences
mazm@amt.ac.cn

ABSTRACT

Graph-based approaches to approximate nearest neighbor search (ANNS) enable fast, high-recall retrieval on billion-scale vector datasets. Among them, the Sparse Neighborhood Graph (SNG) is widely used due to its strong search performance. However, the lack of theoretical understanding of SNG leads to expensive tuning of the truncation parameter that controls graph sparsification. In this work, we present OPT-SNG, a principled framework for analyzing and optimizing SNG construction. We introduce a martingale-based model of the pruning process that characterizes the stochastic evolution of candidate sets during graph construction. Using this framework, we prove that SNG has a maximum out-degree of $O(n^{2/3+\epsilon})$, where $\epsilon > 0$ is an arbitrarily small constant, and an expected search path length of $O(\log n)$. Building on these insights, we derive a closed-form rule for selecting the optimal truncation parameter R , thereby eliminating the need for costly parameter sweeping. Extensive experiments on real-world datasets demonstrate that OPT-SNG achieves an average 5.9× speedup in index construction time, with peak improvements reaching 15.4×, while consistently maintaining or improving search performance.

PVLDB Reference Format:

Xinran Ma, Zhaoqi Zhou, Chuan Zhou, Zaijiu Shang, Guoliang Li, and Zhiming Ma. Sparse Neighborhood Graph-Based Approximate Nearest Neighbor Search Revisited: Theoretical Analysis and Optimization. PVLDB, 14(1): XXX-XXX, 2020.
doi:XX.XX/XXX.XX

PVLDB Artifact Availability:

The source code, data, and/or other artifacts have been made available at <https://github.com/maxila320/OPT-SNG>.

*Both authors contributed equally to this work.

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.

Proceedings of the VLDB Endowment, Vol. 14, No. 1 ISSN 2150-8097.
doi:XX.XX/XXX.XX

1 INTRODUCTION

Nearest Neighbor Search (NNS) is a fundamental problem in computer science [19, 28, 41, 42, 44, 45] and has become a fundamental research topic across various fields such as information retrieval [43, 49, 50], data mining [15, 54, 56], and large language models [4, 35, 66]. Formally, given a dataset $P = \{p_1, \dots, p_n\} \subset \mathbb{R}^d$ of cardinality n and query point $q \in \mathbb{R}^d$, the NNS problem requires finding the k closest points to q under an appropriate distance metric $\|\cdot\|$. The computational complexity of exact NNS grows prohibitively with both dimensionality d due to the well-documented curse of dimensionality [9, 14, 27, 29, 60] and the expanding dataset size n . This fundamental limitation has motivated the development of Approximate Nearest Neighbor Search (ANNS) algorithms, which achieve significant efficiency improvements while sacrificing minimal accuracy.

Modern ANNS approaches achieve remarkable search performance and fall into four categories: tree-based approaches [7, 55], hashing-based approaches [10, 16, 34, 61], vector quantization-based approaches [20, 23], and graph-based approaches [5, 21, 58, 70]. Among these methods, graph-based approaches have demonstrated particular success in practical applications [11, 70], as they effectively capture and exploit local neighborhood structures to enable navigation in high-dimensional spaces. A crucial observation is that despite the diversity of graph construction strategies—which vary in neighbor selection and edge formation rules—most state-of-the-art graph-based ANNS systems share a unified query processing paradigm. The graph-based ANNS framework operates through two complementary phases: (1) an index construction phase that builds a proximity graph preserving neighborhood relationships through various graph refinement strategies, and (2) a query processing phase that employs the *GreedySearch* algorithm [22, 39, 53] to iteratively traverse the graph along decreasing-distance paths until convergence.

Sparse Neighborhood Graph (SNG) is a widely used variant designed to further reduce graph density while preserving practical search quality [3]. It has been widely adopted in several state-of-the-art systems, including DiskANN [53], HNSW [39] and NSG [22].

The worst-case search analysis of SNG by Indyk et al. [30] provides a framework based on aspect ratios¹. Complementary to this, a developing and expected cardinality-based framework is both meaningful and necessary for understanding SNG behavior under random or empirical data distributions.

The construction process of SNG typically involves two main phases: (1) Initial graph generation through a random or approximate k -NN graph, and (2) Graph refinement through iterative pruning[64]. The pruning phase is the computational bottleneck, with $O(n)$ worst-case complexity per node, as it requires scanning all candidate neighbors and evaluating geometric conditions that ensure the graph supports fast convergence of the *GreedySearch* algorithm. To reduce index construction time, a common practice in SNG construction is to impose a maximum node degree via a truncation parameter R . However, existing methods typically determine R through parameter sweeping (or a parameter study) [53], commonly implemented as a binary search guided by the downstream search performance of the constructed graph. This tuning process is computationally expensive, as each iteration requires constructing and evaluating a full index. These limitations highlight the need for a theoretically grounded approach to selecting R —one that ensures high search performance while reducing tuning overhead.

To address both the theoretical and practical challenges in SNG-based approximate nearest neighbor search, we propose **OPT-SNG**, a martingale-based probabilistic framework for analyzing and optimizing SNG construction. In this framework, we characterize the SNG construction process as a discrete stochastic process driven by successive pruning operations. By exploiting the martingale structure naturally arising from the evolution of candidate sets, we establish that the expected trajectory can be characterized by an associated ordinary differential equation (ODE). Furthermore, we prove that the out-degree of each node corresponds precisely to the stopping time of this stochastic process. Within this framework, we derive bounds on two fundamental structural properties of SNGs—the maximum out-degree and the expected search path length. Building upon these theoretical results, we further quantify index construction complexity, which serves as a principled foundation for optimizing the truncation parameter R in SNG-based ANNS systems.

We conclude our contributions as follows.

- We introduce OPT-SNG, the first martingale-based probabilistic model of SNG pruning dynamics. By formulating the pruning process as a submartingale with respect to a natural filtration, we transform graph refinement into an analyzable probabilistic system. This enables the application of concentration inequalities and differential equation methods to study graph evolution.
- We establish a degree bound of $O(n^{2/3+\epsilon})$ for any small constant $\epsilon > 0$, substantially improving over the previous coarse $O(n)$ bound from [53]. This result highlights the inherent sparsity of SNGs and demonstrates that graph structure can be controlled through proper pruning dynamics analysis.
- We prove that GreedySearch on SNGs converges in $O(\log n)$ expected steps. This analytical guarantee validates the intuition

behind SNG’s empirical efficiency and provides theoretical justification.

- Using the above theoretical foundations, we analyze the asymptotic complexity of the truncated SNG construction process and derive a closed-form expression for the optimal truncation parameter R . This yields a principled and automated parameter selection method, avoiding costly parameter sweeping over R . Empirically, our method achieves comparable or superior search performance while reducing index construction time by an average of 5.9× across datasets, with peak speedups reaching 15.4×.

All detailed proofs of lemmas and theorems are provided in the full version [38] of this work on our GitHub repository.

The remainder of this paper is organized as follows. Section 2 reviews the necessary background on SNG-based ANNS, formally defines the problem setting, and introduces the key algorithms considered in this study. In Section 3, we present our probabilistic modeling framework that captures the stochastic dynamics of the pruning process. Sections 4 and 5 derive the main theoretical bounds on graph degree and search path length, respectively. Section 6 describes our optimized construction algorithm and the principled method for truncation parameter selection. Experimental results validating our theoretical predictions and demonstrating practical speedups are reported in Section 7. We survey the related works and position our contributions within the broader literature in Section 8, and Section 9 concludes with a summary and directions for future research.

2 k -ANN PROBLEM AND SPARSE NEIGHBORHOOD GRAPHS

In this section, we formalize the approximate nearest search problem formulation, review existing SNG-based ANNS construction, and summarize the key algorithmic components relevant to our theoretical analysis. Given a dataset $P = \{p_1, \dots, p_n\} \subset \mathbb{R}^d$ of n points in d -dimensional Euclidean space with distance metric $\|\cdot\|$, we consider a query point $q \in \mathbb{R}^d$. See Table 1 for a summary of key notations used throughout.

2.1 Problem Formulation

We formally define the k -approximate nearest neighbor search problem as follows.

Definition 1 (k -Approximate Nearest Neighbor Search). Given a dataset $P \subset \mathbb{R}^d$, a query point $q \in \mathbb{R}^d$, and an integer $k \geq 1$, the k -approximate nearest neighbor search (K-ANNS) problem aims to return a set $A \subseteq P$ of k data points that are closest to q , subject to a desired trade-off between recall and query latency.

Graph-based ANNS methods address this problem by modeling the dataset P as the vertex set of a graph and constructing edges between data points to facilitate efficient search. Typically, the construction process begins with an initial graph, such as a random graph or a k -nearest neighbor graph, and then refines this structure through iterative edge pruning. Formally, these methods construct a graph $G = (P, E)$, where each vertex corresponds to a data point, and the edge set E is designed to preserve neighborhood proximity, enabling fast approximate nearest neighbor search via local graph traversal.

¹The aspect ratio is the ratio of the maximum to minimum distance between any pair of points in the dataset.

Algorithm 1: GreedySearch(G, q, s, k, L)

Input : graph $G = (P, E)$, query point q , entry point s , k for top- k , search width L

Output : top- k approximate nearest neighbors of q

```
1 Visited  $\leftarrow \{s\}$ ;
2 Queue  $\leftarrow$  priority queue initialized with  $s$ ;
3 while  $|Visited| < L$  and  $Queue \neq \emptyset$  do
4    $\hat{p} \leftarrow \text{extract-min}(Queue)$ ;
   /* closest unexpanded point */
5   for each  $p' \in N_{\text{out}}(\hat{p})$  do
6     if  $p' \notin Visited$  then
7       Visited  $\leftarrow Visited \cup \{p'\}$ ;
8       insert  $p'$  into Queue;
9 return the  $k$  points in Visited closest to  $q$ 
```

2.2 GreedySearch Algorithm

Despite differences in index construction, many graph-based ANNS methods share a common query-processing phase: a greedy best-first traversal known as *GreedySearch*. Representative examples include KGraph [18], DiskANN [53], NSG [22], HNSW [39], DPG [36], and SPTAG [13]. The procedure is described in Algorithm 1.

Starting from an entry point s , the algorithm maintains (1) a visited set and (2) a priority queue of candidate nodes ordered by distance to the query q . At each iteration, the closest unexpanded vertex $\hat{p}$ is extracted from the priority queue. Its outgoing neighbors are then explored and added to the priority queue if they have not been visited, and simultaneously, $\hat{p}$ is added to the visited set. The process continues until either the search width reaches the predefined budget L or the priority queue becomes empty. The parameter L controls the search width and governs the recall and latency trade-off: larger values of L explore more candidate nodes, improving recall at the cost of increased search time.

Maintaining and accelerating the convergence of *GreedySearch* is a central objective in graph construction. In particular, ensuring graph connectivity while preserving monotonic search paths is both subtle and essential for guaranteeing fast convergence and high recall. The Sparse Neighborhood Graph (SNG) [3] has been shown to effectively achieve these objectives by iteratively selecting nearest neighbors and pruning redundant candidates [17, 30].

2.3 Sparse Neighborhood Graph Construction

A core operation in SNG construction is the pruning procedure, which selectively retains informative neighbors while aggressively eliminating redundant connections. This procedure is a fundamental building block in several state-of-the-art graph-based ANNS systems, including DiskANN [53], NSG [22], and HNSW [39].

The algorithm proceeds iteratively: in each iteration, it selects the nearest point p^* from the candidate set S , adds p^* to the neighbor set $N_{\text{out}}(p)$, and then prunes candidates from S . Specifically, a candidate p' is removed if it satisfies the blocking condition $\|p - p'\| \geq \alpha \cdot \|p^* - p'\|$. This condition captures the intuition that p' is blocked by the recently selected neighbor p^* —that is, accessing p' through the intermediate neighbor p^* is at least as efficient as a direct connection. Consequently, the direct edge $p \rightarrow p'$

Algorithm 2: SNG-Prune(p, S, α, R)

Input: point p , candidate set S , pruning parameter $\alpha \geq 1$, degree bound R

Output: pruned neighbor set $N_{\text{out}}(p)$ with $|N_{\text{out}}(p)| \leq R$

```
1  $N_{\text{out}}(p) \leftarrow \emptyset$ ;
2 Sort  $S$  in ascending order of  $\|p - \cdot\|$ ;
3 while  $S \neq \emptyset$  do
4    $p^* \leftarrow \arg \min_{p' \in S} \|p - p'\|$ ;
5    $N_{\text{out}}(p) \leftarrow N_{\text{out}}(p) \cup \{p^*\}$ ;
6   Remove  $p^*$  from  $S$ ;
7   if  $|N_{\text{out}}(p)| = R$  then
8     break;
9   for each  $p' \in S$  do
10    if  $\|p - p'\| \geq \alpha \cdot \|p^* - p'\|$  then
11      Remove  $p'$  from  $S$ ;
12 return  $N_{\text{out}}(p)$ ;
```

is redundant and can be safely removed without compromising graph navigability.

Two parameters control the pruning behavior:

Pruning Aggressiveness Parameter α : In the non-relaxed setting ($\alpha = 1$), the blocking condition is symmetric and purely geometric. However, to avoid pathological configurations (e.g., collinear points), the Vamana algorithm [53] introduces $\alpha > 1$, which relaxes the blocking condition and ensures bounded graph diameter. In practice, α is typically fixed in the range $1 \leq \alpha \leq 2$ [31] and does not require extensive tuning, as the SNG structure remains robust across this range.

Degree Truncation Parameter R : To control graph density and construction cost, a maximum degree limit R is imposed. The algorithm terminates after exactly R neighbors have been selected, bounding the out-degree of each node. This truncation ensures that construction remains tractable while allowing a trade-off between graph quality and computational efficiency.

A crucial property of SNG-based construction is that it preserves *monotonic searchability* [22]. Specifically, for any point p , there exists at least one monotonic path consisting of edges that lead to progressively closer points, eventually reaching the nearest neighbor of p . This property is essential for the effectiveness of *GreedySearch*, as it guarantees that each expansion step can make consistent progress toward the target.

Despite its empirical success across multiple state-of-the-art systems (DiskANN, NSG, HNSW), the theoretical understanding of SNG construction remains incomplete. The introduction of relaxation parameter α substantially complicates the analysis: the blocking condition becomes asymmetric, and the geometric properties become more intricate. While Indyk et al. [30] provide worst-case analysis based on aspect ratios and Prokhorenkova and Shekhovtsov [17] establish bounds under specific distributional assumptions, a comprehensive expected-case, cardinality-based framework characterizing SNG behavior has been lacking.

In particular, the truncation parameter R remains a critical tuning knob. Existing practice relies on expensive parameter sweeping: practitioners repeatedly construct indices with different R values, evaluate their search performance through binary or grid search,




Figure 1: Illustration of the SNG process for a single point.

and select the value that optimizes the recall-latency trade-off. This iterative process is computationally prohibitive for large-scale datasets and obscures the fundamental principles governing the relationship between R , graph structure, and search efficiency.

Our OPT-SNG framework addresses these gaps by providing a principled, theoretically grounded approach to analyzing SNG dynamics and optimizing the truncation parameter R .

3 OPT-SNG PROBABILISTIC FRAMEWORK SETUP

Stochastic Modeling of the Pruning Process. To analyze the construction of SNG from a dynamic perspective, we establish the following probabilistic framework that models the evolution of the pruning process. Consider a fixed point $p \in P$ being processed. Let $t \in \{0, 1, 2, \dots\}$ denote the iteration index, where each iteration is associated with two dynamically evolving sets.

Definition 2 (Candidate Set and Processed Set). At iteration t , the *candidate set* is denoted by S_t . It is initialized as $S_0 = P \setminus \{p\}$, meaning that, initially, the point p is allowed to connect to all other points in the dataset. At each iteration, S_t consists of the points that have not yet been pruned and have not been selected as neighbors. The corresponding *processed set* is defined as $S'_t = P \setminus (S_t \cup \{p\})$, which contains all points that have either been pruned or already selected as nearest neighbors in previous iterations.

In the non-truncated setting, the pruning process terminates when $S_t = \emptyset$, or equivalently when $S'_t = P \setminus \{p\}$. Figure 1 provides a schematic illustration of this construction process. Initially ($t = 0$), all other points belong to the candidate set S_0 . At iteration $t = 1$, the nearest neighbor p_1^* is selected, and a subset of points ΔS_1 is pruned and moved into the processed set S'_1 . This process repeats at subsequent iterations: at iteration t , a new nearest neighbor p_t^* is selected from S_{t-1} , additional points are pruned, and the candidate set shrinks accordingly.

Due to the randomness of point distributions in high-dimensional space, the number of pruned points at each iteration is itself random, which induces randomness in the resulting graph structure. Consequently, the pruning procedure can be naturally viewed as a stochastic process. The correspondence between iteration indices and the evolution of S_t and S'_t allows us to formalize the construction using tools from probability theory.

Probability Space for SNG construction. We introduce a probability space $(\Omega, \mathcal{F}, \mathbb{P})$, where Ω denotes the sample space containing all possible sequences of nearest neighbor selections

Table 1: Summary of Notation

Symbol	Description
P	Dataset of n points in $\mathbb{R}^d$
S_t	Candidate set at iteration t
S'_t	Processed set at iteration t : $S'_t = P \setminus (S_t \cup \{p\})$
p_t^*	Nearest neighbor of p in S_t at iteration t
$\mathcal{F}_t$	Natural filtration up to iteration t
ΔS_t	Number of points pruned in iteration t
π_t	Pruning probability at iteration t
R	Truncation parameter (maximum degree limit)
M	Number of processed points at a given level
α	Relaxation parameter
t	Iteration index (time step)

and pruning outcomes during the overall SNG construction. The σ -algebra $\mathcal{F}$ captures all measurable events on Ω , and $\mathbb{P}$ is the associated probability measure. We further define the *natural filtration* $\{\mathcal{F}_t\}_{t \geq 1}$, representing the information available up to iteration t . Specifically,

$$\mathcal{F}_t = \sigma((p_1^*, S_1), (p_2^*, S_2), \dots, (p_t^*, S_t))$$

where $p_i^* = \arg \min_{s \in S_i} \|p - s\|$ is the nearest neighbor selected at iteration i . The candidate set is updated according to the pruning rule $S_{t+1} = S_t \setminus (\{p' \in S_t : \|p - p'\| \geq \alpha \cdot \|p^* - p'\|\} \cup \{p_{t+1}^*\})$ which removes both the selected nearest neighbor and all candidates blocked by it. This formulation captures the pruning process as a sequence of progressively revealed decisions.

Martingale Perspective. We now formalize the dynamic trend of the construction process. Intuitively, at each iteration, the processed set strictly grows while the candidate set shrinks. This monotonic behavior can be rigorously characterized using martingale theory.

Definition 3 (Martingale). A stochastic process $\{X_t\}_{t \geq 0}$ is a *martingale* with respect to a filtration $\{\mathcal{F}_t\}$ if

$$\mathbb{E}[|X_t|] < \infty \quad \text{and} \quad \mathbb{E}[X_{t+1} | \mathcal{F}_t] = X_t$$

for all $t \geq 0$. It is a *supermartingale* if $\mathbb{E}[X_{t+1} | \mathcal{F}_t] \leq X_t$, and a *submartingale* if $\mathbb{E}[X_{t+1} | \mathcal{F}_t] \geq X_t$.

In the context of SNG construction, the size of the candidate set $|S_t|$ forms a supermartingale, while the size of the processed set $|S'_t|$ forms a submartingale with respect to the natural filtration $\{\mathcal{F}_t\}$. This reflects the fact that, in expectation, candidates are eliminated while processed points accumulate over time. We focus on $|S'_t|$ rather than $|S_t|$ because, as complementary sets, $|S'_t|$'s monotonic growth is suited to our analytical framework.

From the SNG pruning rule and the probability framework, we summarize the following properties of processed sets S'_t , which will facilitate the analysis of the construction process. These properties follow directly from our previous results and do not require separate proof.

Lemma 1 (Dynamic Property of S'_t). Given a dataset $P \subset \mathbb{R}^d$ consisting of n points, the SNG construction of $p \in P$ satisfies the properties as follows:

- **Monotonicity:** The processed sets form a nested sequence $\emptyset = S'_0 \subset S'_1 \subset S'_2 \subset \dots \subset S'_T$, with a submartingale $|S'_t|$.
- **Termination:** The final processed set includes $n - 1$ points, excluding only the current point p .
- **Degree property:** Out-degree of p equals the number of iterations, as each iteration adds exactly one directed edge from p to its nearest neighbor in the current candidate set.

Based on this framework, we introduce the following metric to measure the progression of graph construction.

Definition 4 (M - t Level). The SNG construction is said to reach the M - t level at iteration t if t is the smallest iteration that $|S'_t| \geq M$. This represents the *first passage time* at which at least M points have been processed.

The M - t metric serves as a tool to track the progress of graph construction. For a fixed iteration count t , a larger value of M indicates faster pruning and more rapid expansion of the processed set. For a fixed M , a smaller value of t implies that the algorithm has processed M points more quickly.

4 DEGREE BOUNDS OF SNG

In this section, we establish upper bounds on the degree of SNG. As shown in Lemma 1, the out-degree of each point equals the number of pruning iterations during construction, since exactly one edge is added in each iteration. Consequently, bounding the degree reduces to analyzing the dynamics of the pruning process and estimating the total number of iterations required for graph construction.

A preview of the main result. Before delving into the technical analysis, we state the central conclusion of this section: with probability 1, the maximum out-degree of the constructed SNG is $O(n^{2/3+\epsilon})$ for any small $\epsilon > 0$ (Theorem 1, Section 4.5). The remainder of Section 4 formalizes the two-phase pruning dynamics needed for this bound: Lemma 2 (pruning probability), Lemma 3 (2D base case), Lemma 4 (sublinear progress in d dimensions), and finally Theorem 1.

For completeness, all detailed proofs of lemmas and theorems are provided in the full version of this paper on our GitHub repository [1] and arxiv [38].

4.1 Two-Phase Behavior: Empirical Motivation

To motivate our theoretical analysis, we first present empirical evidence revealing a distinctive *two-phase behavior* in SNG construction. Figure 2 illustrates two representative cases from the SIFT1M dataset: a fast case completing in only 31 iterations and a slow case requiring over 1000 iterations. Despite this disparity in total iterations, both cases exhibit a consistent structural pattern: an initial phase of rapid candidate reduction, followed by a prolonged plateau phase during which progress slows significantly.

This empirical observation strongly suggests that a phase-wise analysis is both natural and necessary for deriving degree bounds. In the remainder of this section, we formalize this intuition. Specifically, Lemma 4 establishes rapid progress in the initial phase, which in turn enables a degree bound in Theorem 1.




Figure 2: Evolution of processed and candidate sets during SNG construction. (a) shows a fast pruning case terminating in 31 iterations. (b) shows a slow pruning case requiring over 1,000 iterations.

4.2 Pruning Probability Analysis

We begin by analyzing the pruning probability at each iteration, which governs the rate at which candidate points are eliminated. This is a key ingredient for quantifying the speed of the initial fast pruning phase.

Lemma 2 (Pruning Probability in SNG Construction). *Given a dataset $P \subset \mathbb{R}^d$ consisting of n points independently and uniformly distributed within a ball of radius ρ_0 , centered at p , we consider the SNG construction processed on the center point p with $\alpha = 1$. The probability π_t that a point p' is pruned in the t -th iteration is given by*

$$\pi_t = \frac{I_{1-(\rho_t/2\rho_0)^2} \left(\frac{d+1}{2}, \frac{1}{2} \right) - (\rho_t/\rho_0)^d \cdot I_{0.75} \left(\frac{d+1}{2}, \frac{1}{2} \right)}{2(1 - (\rho_t/\rho_0)^d)},$$

where $I_x(a, b)$ denotes the regularized incomplete Beta function, and ρ_t is the distance between the nearest neighbor and p in t -th iteration. For large d , this probability satisfies

$$\frac{I_{0.75} \left(\frac{d+1}{2}, \frac{1}{2} \right)}{2} < \pi_t < \frac{1}{2}.$$

The complete proof, including the volume computation via the regularized incomplete Beta function and the lower bound for general d , is provided in Appendix B.1. in [38]

4.3 Two-Dimensional Case

To build intuition, we first analyze a low-dimensional setting using the M - t level. In particular, for the two-dimensional case, we show that *only four iterations* suffice to prune over 80% of all candidate points with high probability.

Lemma 3 (Two-Dimensional Case). *Given a dataset $P \subset \mathbb{R}^2$ with n points independently and uniformly distributed within a ball of radius ρ_0 . Then, with probability $1 - o(1)$, the SNG construction process reaches $(0.8(n-1))$ -4 level.*

PROOF. Recall that in the SNG construction of $p \in P$, we iteratively prune a candidate set S_t , starting with $S_0 = P \setminus \{p\}$. The pruning event $A_{p'} = \{\|p - p'\| \geq \alpha\|p_t^* - p'\|\}$ depends only on the relative position of the single candidate p' to the two reference points (p, p_t^*) . Conditional on p_t^* , the remaining candidates are i.i.d.

samples and their relative positions to (p, p_i^*) do not influence one another. Hence the events $\{A_{p'}\}$ are mutually independent given p_i^* , and $\Delta S_i \sim \text{Bin}(|S_{i-1}| - 1, \pi_i)$ for given p_i^* . Let ΔS_i denote the number of points pruned at iteration i . After t iterations, the processed set S'_t consists of all pruned points and the selected nearest neighbors. Therefore, $|S'_t| = \sum_{i=1}^t (\Delta S_i + 1)$, $t = 2, 3, \dots$, with $|S'_1| = \Delta S_1 + 1$, since each iteration contributes exactly one selected nearest neighbor in addition to the pruned points.

From Lemma 2, we know that for $d = 2$, the pruning probability at each iteration satisfies $\pi_i > 1/3$. To estimate the pruning progress, we conservatively assume a uniform lower bound $\pi_{\min} = 1/3$ across all iterations.

The expected number of pruned points at iteration i is $\mathbb{E}[\Delta S_i] = \pi_i(|S_{i-1}| - 1)$, and the variance is $\text{Var}(\Delta S_i) = \pi_i(1 - \pi_i)(|S_{i-1}| - 1)$. Since $\pi_i(1 - \pi_i) < 1/3$, the normalized variance can be bounded as

$$\text{Var}\left(\frac{\Delta S_i}{n-1}\right) = \frac{\pi_i(1 - \pi_i)(|S_{i-1}| - 1)}{(n-1)^2} \leq \frac{1}{3(n-1)}. \quad (1)$$

Under the minimal pruning rate assumption, the expected fraction of points pruned at each iteration is at least

$$\pi_{\min}^1 = \frac{1}{3}, \quad \pi_{\min}^2 = \frac{2}{9}, \quad \pi_{\min}^3 = \frac{4}{27}, \quad \pi_{\min}^4 = \frac{8}{81}.$$

Consequently, the cumulative expected fraction of processed points after four iterations exceeds 0.8.

To ensure that this behavior holds with high probability, we control the deviation of ΔS_i from its expectation by choosing constants $\epsilon_i > 0$ such that

$$\begin{aligned} \epsilon_1 &< \frac{\mathbb{E}[\Delta S_1]}{n-1} - \frac{1}{3}, & \epsilon_2 &< \frac{\mathbb{E}[\Delta S_2]}{n-1} - \frac{2}{9}, \\ \epsilon_3 &< \frac{\mathbb{E}[\Delta S_3]}{n-1} - \frac{4}{27}, & \epsilon_4 &< \frac{\mathbb{E}[\Delta S_4]}{n-1} - \frac{8}{81}. \end{aligned}$$

Applying Chebyshev's inequality yields, for each iteration i ,

$$\begin{aligned} \Pr\left(\left|\frac{\Delta S_i}{n-1} - \frac{\pi_i(|S_{i-1}| - 1)}{n-1}\right| \leq \epsilon_i\right) &\geq 1 - \frac{\text{Var}\left(\frac{\Delta S_i}{n-1}\right)}{\epsilon_i^2} \\ &\geq 1 - \frac{1}{3(n-1)^2 \epsilon_i^2}. \end{aligned} \quad (2)$$

Assuming that all four deviations remain within their respective bounds, the probability that the cumulative number of processed points satisfies $|S'_4| \geq 0.8(n-1)$ is bounded below by

$$\Pr(|S'_4| \geq 0.8(n-1)) \geq \prod_{i=1}^4 \left(1 - \frac{1}{3(n-1)^2 \epsilon_i^2}\right) = 1 - o(1). \quad (4)$$

Therefore, with probability $1 - o(1)$, the SNG construction reaches the $(0.8(n-1))$ -4 level. $\square$

4.4 High-Dimensional Settings

We now extend the analysis to general d -dimensional space. The same phenomenon as in the 2D case manifests as *sublinear-time progress*: a linear number of points is processed within sublinear iterations. This behavior is formalized in the following lemma.

Lemma 4 (Sublinear Time Progress). *Given a dataset $P \subset \mathbb{R}^d$ with n points distributed uniformly, for all relaxation parameter $\alpha > 0$:*

- (1) *For any constants $v_1, v_2 \in (0, 1)$ with $v_1 > v_2$, with probability one, the SNG construction reaches the n^{v_1} - $O(n^{v_2})$ level.*
- (2) *For any constant $v \in (\frac{2}{3}, 1)$, with probability one, the SNG construction reaches the $(n - n^{1-v})$ - $O(n^v)$ level.*

PROOF. Part (1). Fix constants $v_1 > v_2$ and define the first-passage time $t_1 := \inf\{t \geq 0 : |S'_t| \geq n^{v_1}\}$. We prove that $t_1 = O(n^{v_2})$ *almost surely*², i.e., there exists a constant $C > 0$ such that $\Pr(t_1 \leq Cn^{v_2}) = 1$. Equivalently, it suffices to show that for any fixed $C > 0$,

$$\Pr(|S'_{\lfloor Cn^{v_2} \rfloor}| < n^{v_1} \text{ i.o.}) = 0, \quad (5)$$

where "i.o." (infinitely often) denotes the event

$$\limsup_{n \rightarrow \infty} \left\{ |S'_{\lfloor Cn^{v_2} \rfloor}| < n^{v_1} \right\} = \bigcap_{n_0=1}^{\infty} \bigcup_{n \geq n_0} \left\{ |S'_{\lfloor Cn^{v_2} \rfloor}| < n^{v_1} \right\}. \quad (6)$$

Let $t = \lfloor Cn^{v_2} \rfloor$. Since $|S'_t| = \sum_{i=1}^t (\Delta S_i + 1)$, we obtain the deterministic lower bound $|S'_t| \geq t \cdot \min_{1 \leq i \leq t} (\Delta S_i + 1)$, which implies

$$\Pr(|S'_t| < n^{v_1}) \leq \Pr\left(\min_{1 \leq i \leq t} \Delta S_i < \frac{n^{v_1} - t}{t}\right). \quad (7)$$

Since $v_1 > v_2$, for sufficiently large n we have $n^{v_1} > 2t$, and hence

$$\frac{n^{v_1} - t}{t} \geq \frac{1}{2C} n^{v_1 - v_2}. \quad (8)$$

Recall that $\Delta S_i \sim \text{Bin}(|S_{i-1}| - 1, \pi_i)$ with $\mathbb{E}[\Delta S_i] = \pi_i(|S_{i-1}| - 1)$. Moreover, the sampling probability admits the uniform lower bound $\pi_i \geq \frac{1}{2} I_{0.75}\left(\frac{d+1}{2}, \frac{1}{2}\right)$. Applying the Chernoff bound [40], for any $\delta \in (0, 1)$,

$$\begin{aligned} \Pr(\Delta S_i \leq (1 - \delta)\mathbb{E}[\Delta S_i]) &\leq \exp\left(-\frac{\delta^2}{2}\mathbb{E}[\Delta S_i]\right) \\ &\leq \exp\left(-\frac{\delta^2}{4}(|S_{i-1}| - 1) I_{0.75}\left(\frac{d+1}{2}, \frac{1}{2}\right)\right). \end{aligned} \quad (9)$$

For all $i \leq t$ such that $|S_{i-1}| \geq n/2$, we have $\mathbb{E}[\Delta S_i] \geq c_0 n$ for some constant $c_0 > 0$. Choosing $\delta \in (0, 1)$ such that

$$(1 - \delta)\mathbb{E}[\Delta S_i] \geq \frac{1}{2C} n^{v_1 - v_2}, \quad (11)$$

we obtain $\Pr(|S'_t| < n^{v_1}) \leq \exp(-cn)$ for some constant $c > 0$ and all sufficiently large n . Consequently,

$$\sum_{n=1}^{\infty} \Pr(|S'_{\lfloor Cn^{v_2} \rfloor}| < n^{v_1}) < \infty. \quad (12)$$

By the Borel–Cantelli lemma (see Lemma 5 in Appendix A),

$$\Pr(|S'_{\lfloor Cn^{v_2} \rfloor}| < n^{v_1} \text{ i.o.}) = 0, \quad (13)$$

which completes the proof of Part (i).

²Throughout this paper, "almost surely" (abbreviated a.s.) is used in its standard measure-theoretic sense: the event holds with probability 1, i.e., the set of outcomes on which it fails has probability measure zero.

Part (2). Let $\nu \in (2/3, 1)$ and define $t_2 := \inf\{t \geq 0 : |S'_t| \geq n - n^{1-\nu}\}$. We show that $t_2 = O(n^\nu)$ almost surely. Fix $C > 0$ and set $t = \lfloor Cn^\nu \rfloor$. As above,

$$\Pr(|S'_t| < n - n^{1-\nu}) \leq \Pr\left(\min_{1 \leq i \leq t} \Delta S_i < \frac{n - n^{1-\nu} - t}{t}\right). \quad (14)$$

For sufficiently large n , we have $t \leq n/2$ and thus

$$\frac{n - n^{1-\nu} - t}{t} \geq \frac{1}{2C} n^{1-\nu}. \quad (15)$$

Using the same Chernoff bound and the same uniform lower bound on π_i , we again obtain an exponentially small upper bound

$$\Pr(|S'_t| < n - n^{1-\nu}) \leq \exp(-c'n) \quad (16)$$

for some constant $c' > 0$. Therefore,

$$\sum_{n=1}^{\infty} \Pr(|S'_{\lfloor Cn^\nu \rfloor}| < n - n^{1-\nu}) < \infty. \quad (17)$$

Applying the Borel–Cantelli lemma yields

$$\Pr(|S'_{\lfloor Cn^\nu \rfloor}| < n - n^{1-\nu} \text{ i.o.}) = 0, \quad (18)$$

and hence $\Pr(t_2 = O(n^\nu)) = 1$. $\square$

4.5 Main Degree Bound

Lemma 4 establishes that the initial phase rapidly processes a large fraction of points. However, prior empirical observations suggest that this rapid progress does not persist throughout the entire construction. The pruning process undergoes a *significant slowdown* after the initial fast pruning phase, often reaching a *plateau phase*.

To analyze this plateau phase, we exploit a key fact: the submartingale $|S'_t|_{t \geq 0}$ has bounded per-step variance along the pruning trajectory. This enables us to apply the Differential Equation Method (DEM) [33, 59, 62], a standard technique for approximating bounded-variance stochastic processes by deterministic trajectories. DEM yields an ordinary differential equation that captures the expected evolution of $|S'_t|$ in the plateau regime. Combining the initial fast-pruning bound from Lemma 4 with the DEM-based characterization of the plateau phase gives a unified view of the pruning dynamics. We now sketch the proof of our main result, Theorem 1. Detailed proofs of all lemmas and theorems are provided in the full version of this work on our GitHub repository [1] and arxiv [38].

Theorem 1. *Given a dataset P of n points in d -dimensional space, with probability 1, the maximum out-degree in the constructed SNG is bounded by $O(n^{2/3+\epsilon})$ for any small constant $\epsilon > 0$.*

PROOF SKETCH. The proof consists of two parts: (i) establishing a uniform lower bound $\pi_t \geq c_* > 0$ on the one-step pruning probability, and (ii) applying the two-phase pruning analysis from Sections 4.1–4.4, which combines Chernoff bounds, Borel–Cantelli, and Wormald’s differential equation method (DEM). Since the second part only uses the existence of such a positive lower bound, the same argument applies once c_* is verified under each distributional setting.

Case 1: uniform distribution and $\alpha = 1$. For $\alpha = 1$, the blocking condition $\|p - x\| \geq \|p_t^* - x\|$ defines a spherical-cap region. By

Lemma 2, its normalized volume is bounded below by a positive constant: $\pi_t \geq c(d, 1) > 0$. Let $Y_t := |S'_t|$. Since $|S_t| = n - 1 - Y_t$, $Y_{t+1} - Y_t = 1 + \Delta S_{t+1}$, and hence, conditionally on $\mathcal{F}_t$, $\mathbb{E}[Y_{t+1} - Y_t | \mathcal{F}_t] \geq 1 + c(d, 1)(n - 1 - Y_t)$. For $Z_t := Y_t/n$, the normalized drift is therefore governed, up to lower-order terms, by $F(z) = 1 + c(d, 1)(1 - z)$. The DEM conditions hold: boundedness follows from $|Y_{t+1} - Y_t| \leq |S_t| \leq n$, the trend condition follows from the above drift estimate, and F is Lipschitz since

$$|F(z_1) - F(z_2)| = c(d, 1)|z_1 - z_2|. \quad (19)$$

Thus, Wormald’s DEM applies.

By Lemma 4, for any $\nu \in (2/3, 1)$, the first phase reaches $|S'_t| \geq n - n^{1-\nu}$ within $t_1 = O(n^\nu)$ iterations with probability one. At that point at most $n^{1-\nu}$ candidates remain. Applying the DEM-controlled plateau analysis to the residual process gives $O(n^{1-\nu})$ additional iterations. Therefore,

$$T = O(n^\nu) + O(n^{1-\nu}) = O(n^\nu). \quad (20)$$

Taking $\nu = 2/3 + \epsilon$ yields $T = O(n^{2/3+\epsilon})$.

Case 2: uniform distribution and general $\alpha > 0$. For general α , the blocking condition becomes $\|p - x\| \geq \alpha \|p_t^* - x\|$, which defines an Apollonius-type blocking region. Its normalized volume is still bounded below by a positive constant $c(d, \alpha) > 0$, so $\pi_t \geq c(d, \alpha)$. The proof above uses the geometry only through this drift lower bound. Therefore the same argument applies with $c(d, 1)$ replaced by $c(d, \alpha)$, yielding $T = O(n^{2/3+\epsilon})$ for every fixed $\alpha > 0$.

Case 3: bounded non-uniform density. Suppose the local candidate density f satisfies $0 < m \leq f \leq M < \infty$. Conditional on p_t^* , the blocking region $R_t(p_t^*)$ from Case 2 gives

$$\pi_t \geq \frac{m}{M} \frac{\text{Vol}(R_t(p_t^*))}{\text{Vol}(\Omega_{t-1})} \geq \frac{m}{M} c(d, \alpha) =: c_* > 0.$$

Thus non-uniformity only changes the constant in the drift lower bound. The same Chernoff, Borel–Cantelli, and DEM argument then gives $T = O(n^{2/3+\epsilon})$.

Since the out-degree is controlled by the number of pruning iterations, the claimed degree bound follows. Detailed calculations for the α -extension and the bounded-density case are provided in Appendix B.3 in [38]. $\square$

5 SEARCH BOUNDS

We now turn to the analysis of query processing on the constructed SNG. Our goal is to show that, with high probability, *GreedySearch* converges to an approximate nearest neighbor in $O(\log n)$ steps.

Theorem 2. *Given a dataset $P \subset \mathbb{R}^d$ with n points distributed uniformly, with probability 1, the SNG search converges in $O(\log n)$ steps.*

SKETCH PROOF. Let $\{v_0, v_1, \dots, v_k\}$ denote a greedy search path from the start node $p = v_0$ to the approximate nearest neighbor of the query q . At each step, *GreedySearch* moves to the neighbor closest to q , ensuring that $\|v_i - q\| > \|v_{i+1} - q\|$.

To analyze the expected path length k , we analyze the distribution of neighbors around the destination v_k , they are uniformly distributed in a d -dimensional ball of radius ρ_0 . We partition this ball into concentric annular layers and let η_i denote the number of

neighbors lying in the i -th layer. Each η_i follows a binomial distribution proportional to the shell volume, and the total number of neighbors is bounded by $\eta = O(n^{2/3+\epsilon})$ from Theorem 1.

Using volume arguments, we upper bound the expected total number of neighbors located within the ball of radius $\|v_k - v_0\|$ by the expected number of dataset points that fall into the same region. Specifically, we show

$$\mathbb{E}[k \cdot \eta] \leq \frac{n(n+1)^{(d-1)/d} \rho_0^{d+1} \cdot E \left[((\rho_0 - \Delta r)/\rho_0)^{d(k-1)} \right]}{\Delta r^{d+1} \cdot d}, \quad (21)$$

where Δr is the minimum distinguishable pairwise distance difference among all point triples. Applying Jensen's inequality [40], we define an auxiliary function $g(x)$

$$g(x) := \eta \cdot x - \frac{n(n+1)^{(d-1)/d} \rho_0^{d+1}}{d \cdot \Delta r^{d+1}} \left[- \left(\frac{\rho_0 - \Delta r}{\rho_0} \right)^{d(x-1)} + 2 \right], \quad (22)$$

whose root bounds $\mathbb{E}[k]$. We demonstrate $g(k) < 0$, note that $g(x)$ is monotonically increasing, and $g(\log n) > 0$, which together imply that $k < \log n$, hence $\mathbb{E}[k] = O(\log n)$. $\square$

6 TRUNCATION PARAMETER OPTIMIZATION

In the preceding sections, we establish theoretical guarantees for the SNG structure, including sublinear degree growth and logarithmic search complexity. We now leverage these insights to address a practical issue in index construction: how to efficiently select the truncation parameter R , which directly controls the trade-off between construction cost and search performance.

6.1 Complexity Analysis

We summarize the full SNG construction pipeline as implemented in practice, incorporating some widely used optimizations [22, 53]. The process begins by initializing a randomized R -regular graph. For each data point p , a *GreedySearch* procedure is performed from a fixed entry point (e.g., the dataset centroid) to collect an initial candidate set for pruning, with time complexity $T_{\text{search}} = O(R \log n)$. Importantly, in practical implementations the initial candidate set is *not* the entire dataset; instead, it consists only of the points returned by *GreedySearch*. This design significantly reduces construction cost compared to the theoretical baseline that assumes full connectivity. In particular, the size of the initial candidate set satisfies $|S_0| = O(R \log n)$ in expectation. The candidate set is then refined through iterative SNG pruning, where the truncation parameter R bounds the number of pruning iterations. Let $|S_0|$ denote the initial candidate set size; the pruning cost is $T_{\text{prune}} = O(R \cdot |S_0|)$. To further enhance graph connectivity and preserve monotonic searchability, reverse edges are inserted after pruning. These insertions may temporarily violate the degree constraint, triggering additional pruning steps. The amortized cost of reverse-edge handling can be modeled as $T_{\text{reverse}} = O(\alpha \cdot R^3)$, where α is the pruning parameter controlling sparsification aggressiveness.

The computational cost of each stage—*GreedySearch*, SNG pruning, and reverse-edge handling—can be analyzed independently. Aggregating these components yields the overall construction complexity, as expressed below.

$$n(C_1 \cdot R \log n + b_1 + C_2 \cdot (R^2 \log n / \alpha) + b_2 + C_3 \cdot (\alpha R^3) + b_3),$$

Algorithm 3: Optimization of Truncation Parameter R of SNG

Input: test pruning parameter α_1 , target pruning parameter α_2 , number of data points n
Output: optimized truncation parameter R^*

- 1 $R \leftarrow n^{2/3}$; // theoretical upper bound
- 2 Build graph index using $\text{SNG-Prune}(\alpha_1, R)$;
- 3 Compute the average out-degree $\bar{R}$ of the constructed graph;
- 4 $K' \leftarrow \frac{\alpha_1^2 \cdot \bar{R}}{\log n}$;
- 5 $R^* \leftarrow \frac{K' \cdot \log n}{\alpha_2^2}$;
- 6 **return** R^*

where C_i and b_i are implementation-dependent constants.

This expression highlights the competing effects of R and α : increasing R improves graph quality but amplifies both pruning and reverse-edge costs.

6.2 Optimization Principle

As discussed in Section 2.3, α is typically fixed in the range $1 \leq \alpha \leq 2$, whereas the truncation parameter R controls the trade-off between construction cost and search robustness: larger R usually improves recall but increases the cost.

We determine R for a target α using a *marginal optimality principle*. From Section 6.1, the construction cost can be summarized as $n(C_1 R \log n + C_2 R^2 \log n / \alpha + C_3 \alpha R^3 + O(1))$, where the three terms correspond to initial candidate collection, SNG pruning, and reverse-edge handling, respectively. Treating R as a function of α , the dominant stationary condition is

$$-C_2 R^2 \log n / \alpha^2 + C_3 R^3 = 0,$$

which gives $R^* \propto \log n / \alpha^2$. Thus, the optimal truncation parameter grows logarithmically with n and decreases quadratically with α ; the hidden constant absorbs the dimension, implementation constants, and lower-order terms. Algorithm 3 calibrates this constant by a single reference build at $R = n^{2/3}$ under α_1 , measures the average out-degree $\bar{R}$, sets $K' = \alpha_1^2 \bar{R} / \log n$, and returns $R^* = K' \log n / \alpha_2^2$ for the target α_2 . This avoids costly parameter sweeping while using the theoretical degree bound in Theorem 1 as a near-untruncated reference.

R3.D3

7 EXPERIMENTS

We compare a conventional parameter-sweeping pipeline with our *analytical* parameter determination strategy, and assess both construction efficiency and query performance across different graph-based ANNS algorithms and representative datasets.

7.1 Experimental Setup

7.1.1 Datasets. We evaluate our method on six datasets with diverse dimensionalities and application domains, as summarized in Table 2.

The first five datasets are standard benchmarks in the ANNS literature and cover a broad range of real-world applications, including image, audio, and text. **MSong** [8] consists of audio feature vectors for music similarity search. **DEEP1M** [6] and **SIFT1M** [32] are widely used image-retrieval benchmarks based on deep neural

Table 2: Summary of Datasets Used in Experiments

Dataset	Cardinality	Dimension	Query Size	Data Type
MSong	992,272	420	200	Audio
DEEP1M	1,000,000	256	1,000	Image
SIFT1M	1,000,000	128	10,000	Image
GIST1M	1,000,000	960	1,000	Image
GloVe	1,193,514	200	1,000	Text
UNIFORM	1,000,000	128	10,000	Synthetic

features and SIFT descriptors, respectively. **GIST1M** [32] contains high-dimensional GIST descriptors. **GloVe** [47] provides word embeddings for semantic search. In addition, we include a synthetic dataset, **UNIFORM**, where points are generated uniformly at random in a unit hypercube.

7.1.2 Performance Metrics. We measure both index-construction efficiency and query behavior using several indicators. **Index Construction Time** measures the end-to-end time to obtain a finalized graph index, which includes both the time spent on selecting construction parameters and the subsequent construction phase.

For query-time evaluation, **Recall@10** measures the fraction of queries for which at least one true nearest neighbor appears in the top-10 retrieved results, formally defined as $\text{Recall}@k = |A \cap A'|/|A'|$, where A is the retrieved set and A' is the ground truth with $|A| = |A'| = k$. **Queries Per Second (QPS)** measures throughput as the number of queries processed per second and **Query Latency** reports the average time per query in milliseconds at different recall levels.

7.1.3 Parameter Settings. For each fixed α , we vary the search width L (and analogous search-time controls) to generate recall-latency trade-off. For baseline methods that tune the truncation parameter R via parameter sweeping, we employ a binary-search (or grid-search) strategy guided by a normalized 1:1 weighted loss over query recall and latency, which reflects standard practice in iterative tuning pipelines. In contrast, our approach determines R directly from the analytical relationship derived in Section 6.2.

7.1.4 Computing Environment. All experiments are conducted on a workstation equipped with an Intel CPU, 64GB DDR4-3200 memory, and a 2TB NVMe SSD, running Ubuntu 22.04 LTS. We compile all code with GCC 11.4 and enable `-O3` optimizations.

7.1.5 Compared Algorithms. We evaluate three representative SNG-based ANNS algorithms and apply our optimization principle consistently across them. **Vamana** [53], which underlies DiskANN-style indexing, relies on SNG-like pruning and exposes the parameters α and R that directly control graph sparsification and maximum out-degree. **NSG** [22] constructs a navigable graph with a bounded out-degree constraint and typically requires careful tuning of R to balance construction cost and search quality. **HNSW** [39] builds a hierarchical multi-layer graph, where the construction-time candidate list size (e.g., `efConstruction`) plays a similar role in controlling construction overhead and the resulting graph connectivity. We denote the variants using our proposed parameter optimization strategy as **OPT-DiskANN**, **OPT-NSG**, and **OPT-HNSW**, and

Table 3: Index Construction Speedup of OPT-SNG

Dataset	Speedup over Baseline			Avg
	Vamana	NSG	HNSW	
SIFT1M	7.9×	4.8×	3.4×	5.4×
GIST1M	5.6×	7.1×	4.3×	5.7×
DEEP1M	6.4×	6.4×	4.1×	5.6×
MSong	4.2×	5.3×	2.6×	4.0×
GloVe	8.7×	15.4×	2.2×	8.8×
Average	6.6×	7.8×	3.3×	5.9×

refer to their original implementations as **DiskANN**, **NSG**, and **HNSW**, respectively.

7.1.6 Additional baselines. To situate our method within the broader competitive landscape, we additionally compare against four recent state-of-the-art ANNS systems that span graph, hashing, and quantization paradigms: **VSAG** [68], a versatile graph-based ANNS library supporting multiple index structures; **SymphonyQG** [26], a recent graph-quantization method achieving high throughput on standard benchmarks; **DET-LSH** [61], a locality-sensitive hashing method with dynamic encoding trees; and **RaBitQ** [23], a quantization-based method with theoretical error bounds. These methods are not modified by OPT-SNG and they serve as external reference points for query performance.

R3.D1

7.2 Index Construction Performance

We begin by comparing the end-to-end index construction time, including both parameter tuning and the subsequent graph building. Figure 3a reports results across all datasets, where the stacked bars separate tuning time from construction time. Across the board, the most salient difference is the tuning component: parameter sweeping repeatedly triggers expensive partial or full constructions to probe candidate R values, whereas our method largely removes this iterative overhead by computing R analytically once the necessary reference statistics are obtained. As a result, the total construction time decreases substantially even when the actual build phase remains comparable.

Table 3 summarizes the speedups achieved over binary-search baselines. We observe consistent improvements across datasets and algorithms, ranging from 2.2× to 15.4×, with an average speedup of 5.9×. The gains are particularly pronounced on **GloVe**, where NSG achieves a 15.4× speedup, and on **SIFT1M**, where Vamana achieves 7.9×. These large gains occur because the binary-search pipeline typically requires more iterations to converge to a good truncation level on these datasets, leading to large cumulative tuning costs.


7.3 Query Performance Evaluation

We next evaluate whether the analytically chosen parameters (OPT) preserve query-time behavior compared to the parameters selected by parameter sweeping. We report recall and QPS trends by varying the search width L under each α .

Figures 4 present recall versus QPS curves under different α settings. Across datasets, the curves produced by OPT closely track and in several cases slightly improve upon the curves of the original baseline, indicating that removing iterative tuning does not come at the expense of query quality. Intuitively, this behavior is

Table 4: Recall@10 comparison at fixed latency targets. Latency thresholds are dataset- and algorithm-specific, chosen to reflect typical operating points.

Dataset	Vamana		NSG		HNSW	
	Binary	OPT	Binary	OPT	Binary	OPT
MSong	99.26	99.32	99.79	99.82	86.27	91.85
DEEP1M	96.77	98.76	97.94	99.33	98.93	99.71
SIFT1M	99.84	99.90	99.64	99.70	99.92	99.99
GIST1M	99.63	99.76	96.16	96.79	98.20	99.63
GloVe	88.91	89.94	90.34	91.83	92.97	94.61
UNIFORM	98.78	99.62	93.25	95.96	96.53	97.64



(a) Index construction time comparison (b) Ablation: Parameter reuse across datasets

Figure 3: Index construction time comparison and parameter-reuse ablation.

consistent with our derivation: the analytical rule targets the better construction-quality regime that the parameter sweep attempts to locate empirically, but avoids repeatedly paying the cost of probing suboptimal intermediate choices.

To provide a compact comparison under fixed latency budgets, Table 4 summarizes the achieved Recall at representative, dataset and algorithm specific latency targets. The results show that OPT consistently matches or improves Recall@10 across all three algorithms and all datasets. The improvements are most pronounced on DEEP1M, MSong, GIST1M, and GloVe. This slight improvement comes from better calibrating the truncation parameter R with respect to the graph-structure parameter α . When $R \ll R^*$, truncation is too aggressive and removes useful SNG edges, reducing connectivity and recall. When $R \gg R^*$, maintaining R -NN candidates and reverse edges mainly introduces redundant edges and extra overhead, slowing search without further recall gains. The analytically derived R^* can slightly improve Recall by avoiding locally suboptimal choices in discrete parameter sweeps. Notably, HNSW exhibits the largest recall gains on SIFT1M, while on the challenging UNIFORM dataset, the most significant improvement is observed for NSG.

Figure 5 shows that the UNIFORM dataset consistently yields lower throughput than SIFT1M across all evaluated algorithms at comparable recall levels. For example, at 99% Recall@10, Vamana on SIFT1M achieves approximately $7\times$ higher QPS than on UNIFORM. Even under this challenging distribution, OPT consistently achieves higher QPS than the baseline at the same recall levels, demonstrating that our optimization remains effective and preserves query efficiency under uniform data without favorable geometric structure.

Comparison with Additional Baselines. Figure 6 also reports four further baselines, covering the main families of recent ANNS work: two graph-based methods (VSAG and SymphonyQG), a quantization-based method (RaBitQ), a hashing-based method (DET-LSH). Among them, VSAG and SymphonyQG are the most competitive, attaining high QPS at high recall on SIFT1M and Deep1M; nevertheless, OPT-version matches or surpasses their throughput across the full Recall-QPS frontier on every dataset, while requiring substantially less index construction time (on average $1.2\times$ less than VSAG and $1.8\times$ less than SymphonyQG). The non-graph baselines behave very differently: RaBitQ and DET-LSH degrade sharply once recall enters the high-recall regime and never approaching the graph-based frontier. Taken together, the three OPT variants consistently dominate or match every baseline along the entire Recall-QPS frontier.

7.4 Ablation: Parameter Reuse Across Builds

We evaluate a parameter reuse scenario in which the baseline first sweeps R on one dataset and then reuses the resulting value on subsequent builds of the same family, paying no further tuning cost on new data. As shown in Figure 3b, OPT still consistently outperforms the reuse baselines across all five datasets (Sift1M, Gist1M, Deep1M, Msong, GloVe) and three index types (DiskANN, NSG, HNSW). The construction-time reduction averages $1.09\times$ for DiskANN, $1.65\times$ for NSG, and $1.74\times$ for HNSW, and is observed on every (dataset, method) pair. This confirms that the reported construction-time gains stem from improved parameter quality, and that OPT-SNG still provides direct construction speedups even after removing the effect of parameter-sweep overhead.

7.5 Effect of Pruning Parameter α

Figure 4 illustrates the recall versus QPS curves under different α settings. Across almost all datasets, increasing the pruning parameter α exhibits a consistent trend: for a fixed recall level, the achievable QPS initially increases as α grows, and then either saturates or slightly degrades when α becomes too large. We observe that $\alpha = 1.50$ provides the best overall balance across datasets, achieving the highest or near-highest QPS at comparable recall levels in most cases.

The sensitivity to α varies across datasets. High-dimensional datasets such as GIST1M and DEEP1M exhibit more pronounced performance degradation at large α values, as overly sparse graphs struggle to preserve navigability in high-dimensional spaces. In contrast, lower- or medium-dimensional datasets such as SIFT1M, GloVe, and MSong remain robust across a wider range of α values, showing only mild performance variation.

7.6 Scalability Study

Finally, we evaluate the scalability of our approach by varying the dataset size and reporting the index construction time. In addition, we conduct experiments on the SIFT10M dataset by sampling subsets of size 2M, 4M, 6M, and 8M points. As shown in Figure 7, the construction speedups achieved by the OPT variants are consistently preserved across all sampled subsets, and the benefits become increasingly pronounced as the dataset size grows.

Furthermore, as illustrated in Figure 8, the recall-QPS curves produced by OPT consistently match or outperform those obtained




Figure 4: Query performance across varying α values on real-world datasets




Figure 5: Query performance across varying α values on UNIFORM

via parameter sweeping across all three algorithms, without compromising search quality.

To verify scalability beyond the million-scale setting, we evaluate OPT-Vamana on three standard billion-scale benchmarks: **SIFT1B** (1B vectors, 128-dim), **DEEP1B** (1B vectors, 96-dim), and **SpaceV1B** (1B vectors, 100-dim), following the experimental protocol used by recent large-scale ANNS evaluations such as [63]. All baselines and OPT variants are run on the same hardware described in Section 7, and we report results in Figure 9.

Across all three datasets, OPT-SNG preserves the construction-time advantage observed at the million-point scale. In particular, it achieves graph-construction speedups of **2.48×** on SIFT1B, **1.03×** on DEEP1B, and **1.07×** on SpaceV1B. Moreover, at fixed latency

targets, OPT-SNG attains Recall@10 higher than that of the original version on all three datasets.

7.7 Summary

The experimental results provide evidence supporting the effectiveness of the proposed analytical parameter optimization strategy. Across all evaluated datasets and SNG-based algorithms, our method substantially reduces index construction time by eliminating expensive iterative parameter sweeping. In particular, tuning overhead is reduced by more than an order of magnitude, leading to overall construction speedups of up to 15.4× and an average speedup of 5.9× across datasets.

R1.05,
R3.W3




Figure 6: Comparison with additional baselines on six datasets




Figure 7: Scalability of index construction speedup




Figure 8: Query performance on SIFT10M




Figure 9: Query performance on billion datasets

Crucially, this acceleration does not come at the cost of query quality. The analytically determined parameters consistently preserve—and in many cases improve the recall and QPS compared to baseline parameter sweeping. Recall is matched or improved across all datasets and algorithms at representative latency targets.

8 RELATED WORK

Approximate Nearest Neighbor Search. ANNS has evolved to support diverse scenarios, including multi-metric spaces [69], hybrid queries [46], privacy-preserving federated settings [67], and retrieval-augmented generation [35]. Existing methods are commonly categorized into tree-based [55, 57], hashing-based [16, 34, 61], quantization-based [20], and graph-based approaches [37, 58, 70]. Tree-based methods, such as k -d trees [7], hierarchically partition the space but suffer from the curse of dimensionality. Hashing-based methods, such as LSH [16], map similar points to the same bucket with high probability but often incur substantial memory overhead. Quantization-based methods, including product quantization (PQ) [32] and optimized product quantization (OPQ) [24], compress vectors for efficient distance computation at the cost of precision loss. Hybrid systems combining graph search with quantization, such as DiskANN [53] and SPANN [11], further support billion-scale search.

Graph-Based ANNS Framework. Graph-based ANNS systems construct sparse proximity graphs to balance search efficiency and index size. Representative systems include HNSW [39], NSG [22], KGraph [18], DiskANN [53], DPG [36], and SPTAG [13]. Extensions such as Filtered-DiskANN [25] and FreshDiskANN [51] adapt this

framework to filtered or dynamic large-scale settings, while recent work explores cross-modal retrieval [12] and filtered search [2, 46].

Existing theoretical studies mainly analyze structural properties such as degree bounds, query complexity, and search path length under simplified distributional assumptions. Laarhoven [48] establishes sublinear query complexity for threshold-based nearest neighbor graphs under spherical data models, where sufficient local connectivity enables efficient greedy routing. Prokhorenkova and Shekhovtsov [17] analyze navigable graphs under a binomial distribution and show that average degree $O(n^\beta)$ for $\beta < 1/2$ can preserve navigability, revealing a trade-off between graph sparsity and search reliability.

Parameter Tuning for ANNS Systems. Selecting optimal parameters for ANNS systems remains challenging. Traditional parameter sweeping requires repeated index construction and evaluation, which is computationally expensive. Automated approaches such as VDTuner [65] use learning-based models to predict configurations. Our work takes a complementary route by deriving analytical parameter-selection rules from the pruning dynamics, thereby reducing the need for iterative search.

9 CONCLUSIONS

This paper revisits SNG-based graph indexing for approximate nearest neighbor search from theoretical and optimization perspectives. We develop OPT-SNG, a martingale-based probabilistic framework for modeling the stochastic pruning dynamics in SNG construction. This framework enables a principled analysis of graph structures generated by widely used systems. We show that the constructed SNG has maximum out-degree $O(n^{2/3+\epsilon})$ for any $\epsilon > 0$, and that *GreedySearch* converges in an expected $O(\log n)$ steps, giving an expected-case, cardinality-based complement to existing analyses.

Building on these guarantees, we derive a closed-form rule for selecting the truncation parameter R as a function of α and n , together with a lightweight calibration procedure for estimating the scaling constant. Experiments on real-world benchmarks and a synthetic UNIFORM dataset show that our analytical parameter selection removes expensive parameter sweeping, reducing construction time by 5.9× on average while maintaining or improving search performance.

An important direction for future work is to extend Theorem 2 beyond the current uniform-distribution setting. A rigorous extension would require re-establishing the volume-ratio inequalities and drift conditions used in the proof under non-uniform distributions.

R1.03

REFERENCES

- [1] [n.d.]. OPT-SNG: Sparse Neighborhood Graph-Based Approximate Nearest Neighbor Search Revisited. <https://github.com/maxila320/OPT-SNG>. GitHub repository.
- [2] Anas Ait Aomar, Karima Echihabi, Marco Arnaboldi, Ioannis Alagiannis, Damien Hilloulin, and Manal Cherkaoui. 2025. RWalks: Random Walks as Attribute Diffusers for Filtered Vector Search. *Proceedings of the ACM on Management of Data* 3, 3 (2025), 1–26.
- [3] Sunil Arya and David M. Mount. 1993. Approximate nearest neighbor queries in fixed dimensions. In *Proceedings of the 4th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*. SIAM, Philadelphia, PA, USA, 271–280.
- [4] Akari Asai, Sewon Min, Zexuan Zhong, and Danqi Chen. 2023. Retrieval-based Language Models and Applications. *ACL Tutorial* (2023).
- [5] Ilias Azizi, Karima Echihabi, and Themis Palpanas. 2023. Elpis: Graph-Based Similarity Search for Scalable Data Science. *Proc. VLDB Endow.* 16, 6 (2023), 1548–1559.
- [6] Artem Babenko and Victor S. Lempitsky. 2016. Efficient Indexing of Billion-Scale Datasets of Deep Descriptors. In *2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016, Las Vegas, NV, USA, June 27–30, 2016*. IEEE Computer Society, 2055–2063.
- [7] Jon Louis Bentley. 1975. Multidimensional Binary Search Trees Used for Associative Searching. *Commun. ACM* 18, 9 (1975), 509–517.
- [8] Thierry Bertin-Mahieux, Daniel P. W. Ellis, Brian Whitman, and Paul Lamere. 2011. The Million Song Dataset. In *Proceedings of the 12th International Society for Music Information Retrieval Conference, ISMIR 2011*. University of Miami, 591–596.
- [9] Allan Borodin, Rafail Ostrovsky, and Yuval Rabani. 1999. Lower bounds for high dimensional nearest neighbor search and related problems. In *Proceedings of the thirty-first annual ACM symposium on Theory of computing*. 312–321.
- [10] Deng Cai. 2021. A Revisit of Hashing Algorithms for Approximate Nearest Neighbor Search. *IEEE Trans. Knowl. Data Eng.* 33, 6 (2021), 2337–2348.
- [11] Meng Chen, Kai Zhang, Zhenying He, Yanan Jing, and X. Sean Wang. 2024. RoarGraph: A Projected Bipartite Graph for Efficient Cross-Modal Approximate Nearest Neighbor Search. *Proc. VLDB Endow.* 17, 11 (2024), 2735–2749.
- [12] Meng Chen, Kai Zhang, Zhenying He, Yanan Jing, and X Sean Wang. 2024. Roargraph: A projected bipartite graph for efficient cross-modal approximate nearest neighbor search. *arXiv preprint arXiv:2408.08933* (2024).
- [13] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. 2021. Spann: Highly-efficient billion-scale approximate nearest neighborhood search. *Advances in Neural Information Processing Systems* 34 (2021), 5199–5212.
- [14] Romain Couillet and Zhenyu Liao. 2022. *Random Matrix Methods for Machine Learning*. Cambridge University Press.
- [15] Thomas Cover and Peter Hart. 1967. Nearest neighbor pattern classification. *IEEE transactions on information theory* 13, 1 (1967), 21–27.
- [16] Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S. Mirrokni. 2004. Locality-sensitive hashing scheme based on p-stable distributions. In *Proceedings of the 20th ACM Symposium on Computational Geometry, Brooklyn, New York, USA, June 8–11, 2004*. ACM, 253–262.
- [17] Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, and Torsten Suel. 2024. Navigable Graphs for High-Dimensional Nearest Neighbor Search: Constructions and Limits. In *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*.
- [18] Wei Dong, Charikar Moses, and Kai Li. 2011. Efficient k-nearest neighbor graph construction for generic similarity measures. In *Proceedings of the 20th international conference on World wide web*. 577–586.
- [19] Karima Echihabi, Kostas Zoumpatianos, and Themis Palpanas. 2020. Scalable Machine Learning on High-Dimensional Vectors: From Data Series to Deep Network Embeddings. In *International Conference on Web Intelligence, Mining and Semantics WIMS*. 1–6.
- [20] Jingya Fan, Yang Wang, Wenwen Song, and Zhibin Pan. 2024. Flexible product quantization for fast approximate nearest neighbor search. *Multim. Tools Appl.* 83, 18 (2024), 53243–53261.
- [21] Cong Fu, Changxu Wang, and Deng Cai. 2022. High Dimensional Similarity Search With Satellite System Graph: Efficiency, Scalability, and Unindexed Query Compatibility. *IEEE Trans. Pattern Anal. Mach. Intell.* 44, 8 (2022), 4139–4150.
- [22] Cong Fu, Chao Xiang, Changxu Wang, and Deng Cai. 2019. Fast Approximate Nearest Neighbor Search With The Navigating Spreading-out Graph. *Proc. VLDB Endow.* 12, 5 (2019), 461–474.
- [23] Jianyang Gao and Cheng Long. 2024. RaBitQ: Quantizing High-Dimensional Vectors with a Theoretical Error Bound for Approximate Nearest Neighbor Search. *Proc. ACM Manag. Data* 2, 3 (2024), 167.
- [24] Tiezheng Ge, Kaiming He, Qifa Ke, and Jian Sun. 2013. Optimized product quantization. *IEEE transactions on pattern analysis and machine intelligence* 36, 4 (2013), 744–755.
- [25] Siddharth Gollapudi, Neel Karia, Varun Sivashankar, Ravishankar Krishnaswamy, Nikit Begwani, Swapnil Raz, Yiyong Lin, Yin Zhang, Neelam Mahapatro, Premkumar Srinivasan, et al. 2023. Filtered-diskann: Graph algorithms for approximate nearest neighbor search with filters. In *Proceedings of the ACM Web Conference 2023*. 3406–3416.
- [26] Yutong Gou, Jianyang Gao, Yuxuan Xu, and Cheng Long. 2025. SymphonyQG: towards symphonious integration of quantization and graph for approximate nearest neighbor search. *Proceedings of the ACM on Management of Data* 3, 1 (2025), 1–26.
- [27] Alexander Hinneburg, Charu C Aggarwal, and Daniel A Keim. 2000. What is the nearest neighbor in high dimensional spaces?. In *26th Internat. Conference on Very Large Databases*. 506–515.
- [28] Qiang Huang and Anthony Kum Hoe Tung. 2023. Lightweight-Yet-Efficient: Revitalizing Ball-Tree for Point-to-Hyperplane Nearest Neighbor Search. In *39th IEEE International Conference on Data Engineering, ICDE 2023, Anaheim, CA, USA, April 3–7, 2023*. IEEE, 436–449.
- [29] Piotr Indyk and Rajeev Motwani. 1998. Approximate nearest neighbors: towards removing the curse of dimensionality. In *Proceedings of the thirtieth annual ACM symposium on Theory of computing*. 604–613.
- [30] Piotr Indyk and Haikun Xu. 2023. Worst-case performance of popular approximate nearest neighbor search implementations: Guarantees and limitations. *Advances in Neural Information Processing Systems* 36 (2023), 66239–66256.
- [31] Shikhar Jaiswal, Ravishankar Krishnaswamy, Ankit Garg, Harsha Vardhan Simhadri, and Sheshansh Agrawal. 2022. Ood-diskann: Efficient and scalable graph anns for out-of-distribution queries. *arXiv preprint arXiv:2211.12850* (2022).
- [32] Hervé Jégou, Matthijs Douze, and Cordelia Schmid. 2011. Product Quantization for Nearest Neighbor Search. *IEEE Trans. Pattern Anal. Mach. Intell.* 33, 1 (2011), 117–128.
- [33] Joseph Lauer and Nicholas Wormald. 2007. Large independent sets in regular graphs of large girth. *Journal of Combinatorial Theory, Series B* 97, 6 (2007), 999–1009.
- [34] Yifan Lei, Qiang Huang, Mohan S. Kankanhalli, and Anthony K. H. Tung. 2020. Locality-Sensitive Hashing Scheme based on Longest Circular Co-Substring. In *Proceedings of the 2020 International Conference on Management of Data, SIGMOD Conference 2020, online conference [Portland, OR, USA], June 14–19, 2020*. ACM, 2589–2599.
- [35] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6–12, 2020, virtual*.
- [36] Wen Li, Ying Zhang, Yifang Sun, Wei Wang, Mingjie Li, Wenjie Zhang, and Xuemin Lin. 2019. Approximate nearest neighbor search on high dimensional data—experiments, analyses, and improvement. *IEEE Transactions on Knowledge and Data Engineering* 32, 8 (2019), 1475–1488.
- [37] Kejing Lu, Chuan Xiao, and Yoshiharu Ishikawa. 2024. Probabilistic Routing for Graph-Based Approximate Nearest Neighbor Search. In *Forty-first International Conference on Machine Learning, ICML 2024, Vienna, Austria, July 21–27, 2024*. OpenReview.net.
- [38] Xinran Ma, Zhaoqi Zhou, Chuan Zhou, Qi Meng, Zaijiu Shang, Guoliang Li, and Zhiming Ma. 2025. Graph-Based Approximate Nearest Neighbor Search Revisited: Theoretical Analysis and Optimization. *arXiv preprint arXiv:2509.15531* (2025). arXiv:2509.15531 [cs.DS]
- [39] Yury A. Malkov and Dmitry A. Yashunin. 2020. Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. *IEEE Trans. Pattern Anal. Mach. Intell.* 42, 4 (2020), 824–836.
- [40] Michael Mitzenmacher and Eli Upfal. 2017. *Probability and computing: Randomization and probabilistic techniques in algorithms and data analysis*. Cambridge university press. 66–75 pages.
- [41] Cun Matthew Mu, Jun Raymond Zhao, Guang Yang, Binwei Yang, and Zheng John Yan. 2019. Fast and Exact Nearest Neighbor Search in Hamming Space on Full-Text Search Engines. In *Similarity Search and Applications - 12th International Conference, SISAP 2019, Newark, NJ, USA, October 2–4, 2019, Proceedings (Lecture Notes in Computer Science)*, Vol. 11807. Springer, 49–56.
- [42] Riley Murray, James Demmel, Michael W. Mahoney, N. Benjamin Erichson, Maksim Melnichenko, Osman Asif Malik, Laura Grigori, Piotr Luszczek, Michał Dereziński, Miles E. Lopes, Tianyu Liang, Hengrui Luo, and Jack Dongarra. 2023. Randomized Numerical Linear Algebra : A Perspective on the Field With an Eye to Software. arXiv:2302.11474
- [43] Shumpei Okura, Yukihiro Tagami, Shingo Ono, and Akira Tajima. 2017. Embedding-based News Recommendation for Millions of Users. In *SIGKDD*. ACM, 1933–1942.
- [44] Themis Palpanas. 2015. Data Series Management: The Road to Big Sequence Analytics. *SIGMOD Record* (2015).
- [45] Themis Palpanas and Volker Beckmann. 48(3), 2019. Report on the First and Second Interdisciplinary Time Series Analysis Workshop (ITISA). *SIGREC* (48(3), 2019).

- [46] Liana Patel, Peter Kraft, Carlos Guestrin, and Matei Zaharia. 2024. ACORN: Performant and Predicate-Agnostic Search Over Vector Embeddings and Structured Data. *Proc. ACM Manag. Data* 2, 3 (2024), 120.
- [47] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. Glove: Global Vectors for Word Representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing, EMNLP 2014, A meeting of SIGDAT, a Special Interest Group of the ACL*. ACL, 1532–1543.
- [48] Liudmila Prokhorenkova and Aleksandr Shekhovtsov. 2020. Graph-based Nearest Neighbor Search: From Practice to Theory. In *Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13-18 July 2020, Virtual Event (Proceedings of Machine Learning Research)*, Vol. 119. PMLR, 7803–7813.
- [49] M. M. Mahabubur Rahman and Jelena Tesic. 2022. Evaluating Hybrid Approximate Nearest Neighbor Indexing and Search (HANNIS) for High-dimensional Image Feature Search. In *IEEE International Conference on Big Data, Big Data 2022, Osaka, Japan, December 17-20, 2022*. IEEE, 6802–6804.
- [50] Susav Shrestha, Narasimha Reddy, and Zongwang Li. 2024. ESPN: Memory-Efficient Multi-vector Information Retrieval. In *Proceedings of the 2024 ACM SIGPLAN International Symposium on Memory Management, ISMM 2024, Copenhagen, Denmark, 25 June 2024*. ACM, 95–107.
- [51] Aditi Singh, Suhas Jayaram Subramanya, Ravishankar Krishnaswamy, and Harsha Vardhan Simhadri. 2021. FreshDiskANN: A Fast and Accurate Graph-Based ANN Index for Streaming Similarity Search. *CoRR* abs/2105.09613 (2021). arXiv:2105.09613 <https://arxiv.org/abs/2105.09613>
- [52] Sunil Srinivasa and Martin Haenggi. 2010. Distance Distributions in Finite Uniformly Random Networks: Theory and Applications. *IEEE Trans. Veh. Technol.* 59, 2 (2010), 940–949.
- [53] Suhas Jayaram Subramanya, Devvrit Fnu, Harsha Vardhan Simhadri, Ravishankar Krishnaswamy, and Rohan Kadekodi. 2019. DiskANN: Fast Accurate Billion-point Nearest Neighbor Search on a Single Node. In *Advances in Neural Information Processing Systems*, H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché Buc, E. Fox, and R. Garnett (Eds.), Vol. 32. Curran Associates, Inc.
- [54] Yukihiko Tagami. 2017. AnnexML: Approximate Nearest Neighbor Search for Extreme Multi-label Classification. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, Halifax, NS, Canada, August 13 - 17, 2017*. ACM, 455–464.
- [55] Yufei Tao, Ke Yi, Cheng Sheng, and Panos Kalnis. 2009. Quality and efficiency in high dimensional nearest neighbor search. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2009, Providence, Rhode Island, USA, June 29 - July 2, 2009*. ACM, 563–576.
- [56] Hao Wang, Tong Xu, Qi Liu, Defu Lian, Enhong Chen, Dongfang Du, Han Wu, and Wen Su. 2019. MCNE: An end-to-end framework for learning multiple conditional network representations of social network. In *Proceedings of the 25th ACM SIGKDD international conference on knowledge discovery & data mining*. 1064–1072.
- [57] Jingdong Wang, Naiyan Wang, You Jia, Jian Li, Gang Zeng, Hongbin Zha, and Xian-Sheng Hua. 2014. Trinary-Projection Trees for Approximate Nearest Neighbor Search. *IEEE Trans. Pattern Anal. Mach. Intell.* 36, 2 (2014), 388–403.
- [58] Mengzhao Wang, Lingwei Lv, Xiaoliang Xu, Yuxiang Wang, Qiang Yue, and Jiongkang Ni. 2023. An Efficient and Robust Framework for Approximate Nearest Neighbor Search with Attribute Constraint. In *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10 - 16, 2023*.
- [59] Lutz Warnke. 2019. On Wormald's differential equation method. *arXiv preprint arXiv:1905.08928* (2019).
- [60] Roger Weber, Hans-Jörg Schek, and Stephen Blott. 1998. A quantitative analysis and performance study for similarity-search methods in high-dimensional spaces. In *VLDB*, Vol. 98. 194–205.
- [61] Jiuqi Wei, Botao Peng, Xiaodong Lee, and Themis Palpanas. 2024. DET-LSH: A Locality-Sensitive Hashing Scheme with Dynamic Encoding Tree for Approximate Nearest Neighbor Search. *Proc. VLDB Endow.* 17, 9 (2024), 2241–2254.
- [62] Nicholas Wormald. 1995. Differential equations for random processes and random graphs. *The annals of applied probability* (1995), 1217–1235.
- [63] Qian Xu, Juan Yang, Feng Zhang, Junda Pan, Kang Chen, Youren Shen, Amelie Chi Zhou, and Xiaoyong Du. 2025. Tribase: A vector data query engine for reliable and lossless pruning compression using triangle inequalities. *Proceedings of the ACM on Management of Data* 3, 1 (2025), 1–28.
- [64] Shuo Yang, Jiadong Xie, Yingfan Liu, Jeffrey Xu Yu, Xiyue Gao, Qianru Wang, Yanguo Peng, and Jiangtao Cui. 2025. Revisiting the Index Construction of Proximity Graph-Based Approximate Nearest Neighbor Search. *Proc. VLDB Endow.* 18, 6 (2025), 1825–1838.
- [65] Tiannuo Yang, Wen Hu, Wangqi Peng, Yusen Li, Jianguo Li, Gang Wang, and Xiaoguang Liu. 2024. Vdtuner: Automated performance tuning for vector data management systems. In *2024 IEEE 40th International Conference on Data Engineering (ICDE)*. IEEE, 4357–4369.
- [66] Qiang Yue, Xiaoliang Xu, Yuxiang Wang, Yikun Tao, and Xuliyan Luo. 2024. Routing-Guided Learned Product Quantization for Graph-Based Approximate Nearest Neighbor Search. In *40th IEEE International Conference on Data Engineering, ICDE 2024, Utrecht, The Netherlands, May 13-16, 2024*. IEEE, 4870–4883.
- [67] Xinyi Zhang, Qichen Wang, Cheng Xu, Yun Peng, and Jianliang Xu. 2024. Fed-KNN: Secure Federated k-Nearest Neighbor Search. *Proc. ACM Manag. Data* 2, 1 (2024), V2mod011:1–V2mod011:26.
- [68] Xiaoyao Zhong, Haotian Li, Jiabao Jin, Mingyu Yang, Deming Chu, Xiangyu Wang, Zhitao Shen, Wei Jia, George Gu, Yi Xie, Xuemin Lin, Heng Tao Shen, Jingkuan Song, and Peng Cheng. 2025. VSAG: An Optimized Search Framework for Graph-based Approximate Nearest Neighbor Search. *Proc. VLDB Endow.* 18, 12 (2025), 5017–5030.
- [69] Yifan Zhu, Lu Chen, Yunjun Gao, Ruiyao Ma, Baihua Zheng, and Jingwen Zhao. 2024. HJG: An Effective Hierarchical Joint Graph for ANNS in Multi-Metric Spaces. In *40th IEEE International Conference on Data Engineering, ICDE 2024, Utrecht, The Netherlands, May 13-16, 2024*. IEEE, 4275–4287.
- [70] Chaoji Zuo, Miao Qiao, Wenchao Zhou, Feifei Li, and Dong Deng. 2024. SeRF: Segment Graph for Range-Filtering Approximate Nearest Neighbor Search. *Proc. ACM Manag. Data* 2, 1 (2024), 69:1–69:26.